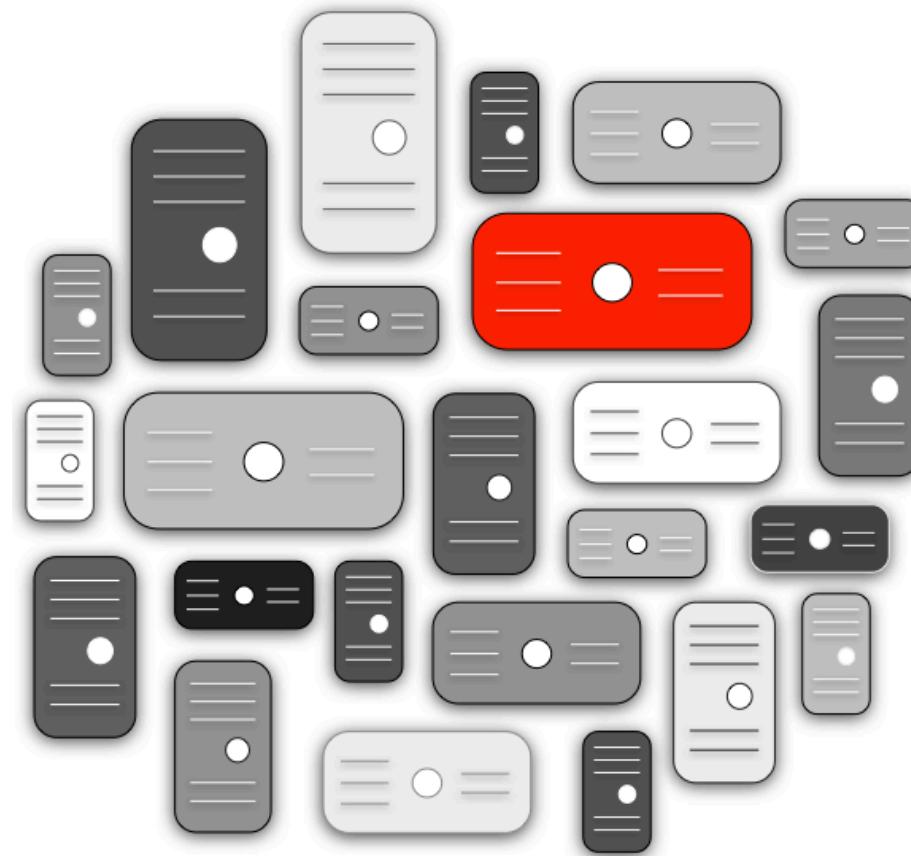


# HotSpot: Automated Server Hopping in Cloud Spot Markets



Supreeth Shastri and David Irwin

UMassAmherst

# Transient Servers are Ubiquitous in the Cloud

Servers that may **terminate anytime** after an **advance warning** period



**Spot Instances:** variable-priced transient VMs offered via second price auction



**Preemptible VM:** short-lived VMs offered at fixed but discounted prices



**Internal Use:** Resource harvesting in datacenters  
[SoCC 2016. OSDI 2016, ATC 2017]

# EC2 Spot Markets in a Nutshell

7600+ spot markets worldwide



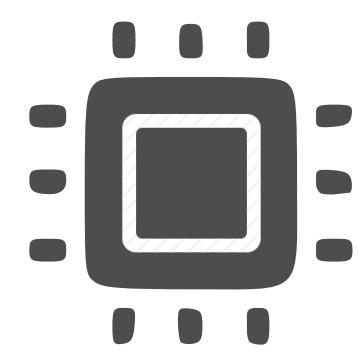
1

The users bid for  
VMs in a second  
price auction



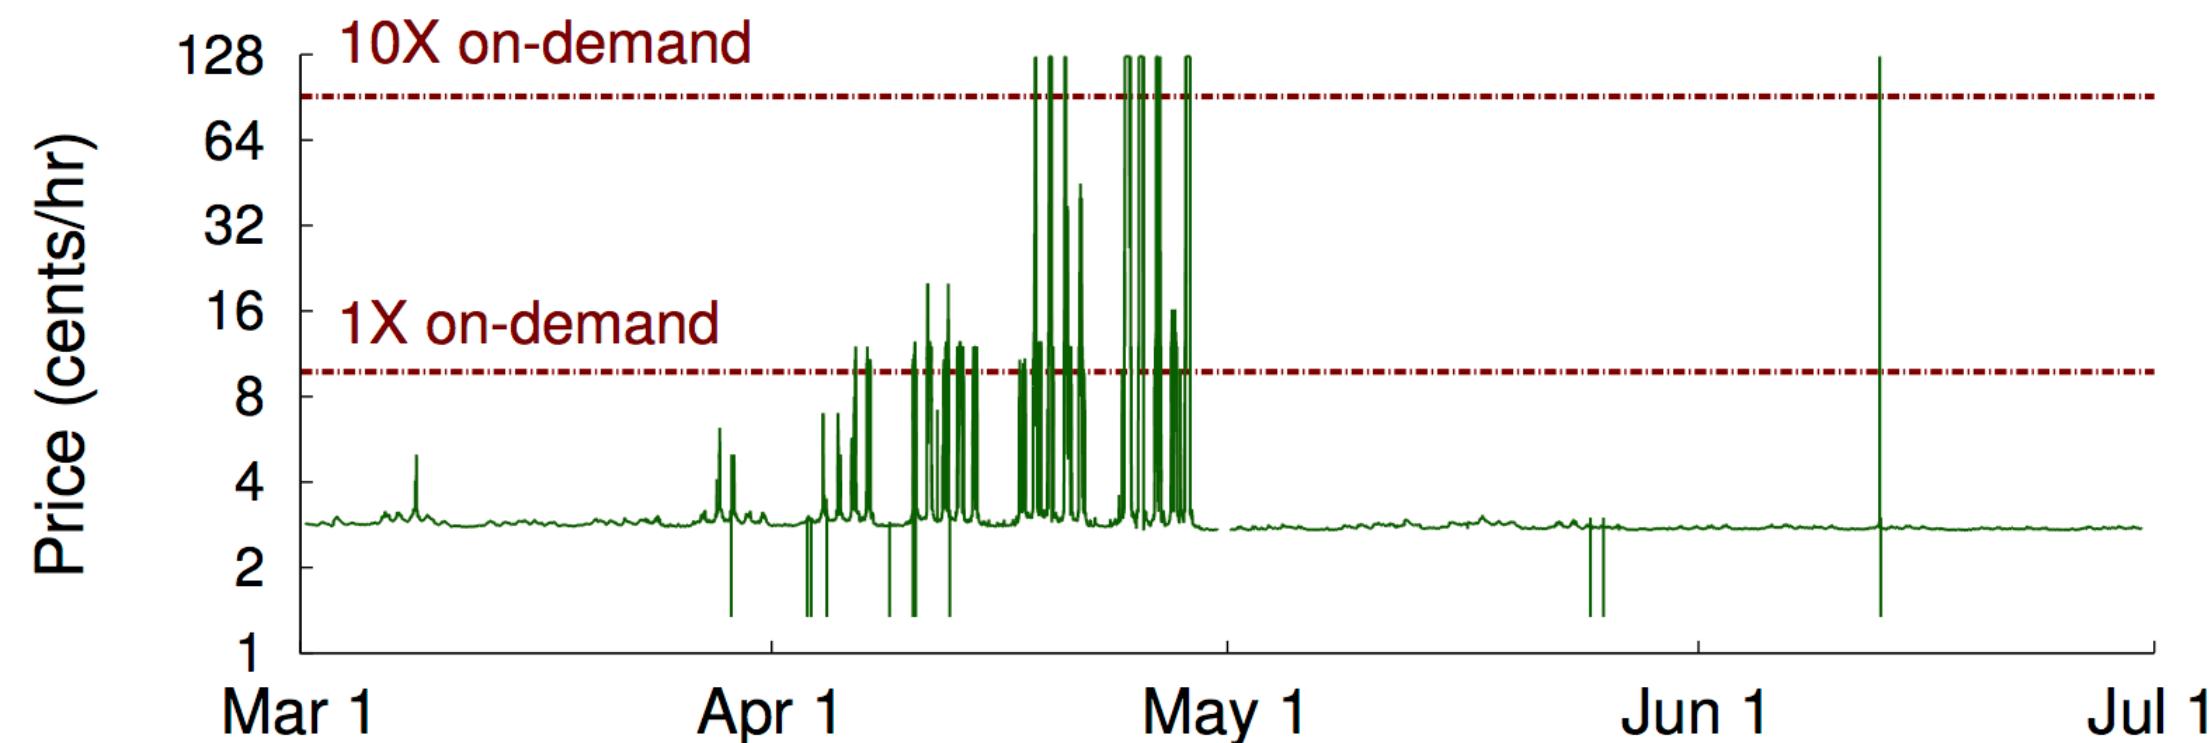
2

EC2 evaluates supply-  
demand dynamic to  
price spot servers



3

EC2 allocates if bid  
price  $\geq$  spot price;  
Revokes when not.



The defining characteristics of spot VMs are  
*low average price* and *unexpected revocations*

Applications and frameworks do not perform well  
when the underlying servers are frequently revoked

Prior work treats ***revocations as failures***,  
and employs fault-tolerance to reduce its impact

... but  
insurance-like approaches ignore

2015

SpotOn [**SoCC**]

SpotCheck [**EuroSys**]

Cumulon [**VLDB**]

2016

TR-Spark [**SoCC**]

Flint [**EuroSys**]

BOSS [**Infocom**]

2017

Proteus [**EuroSys**]

Pado [**EuroSys**]

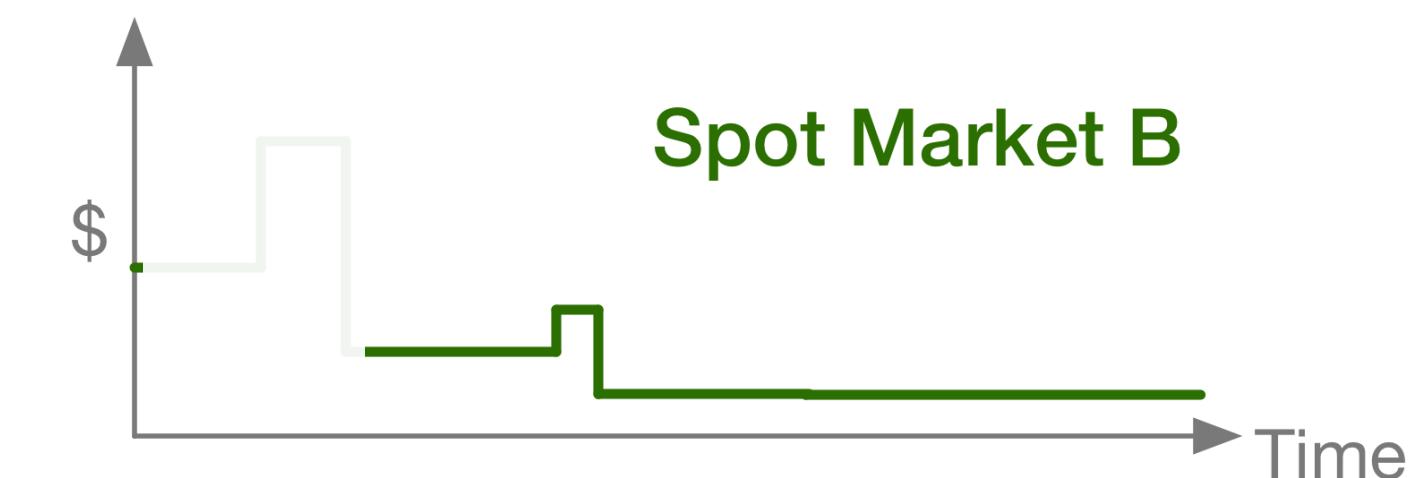
Exosphere [**Sigmetrics**]

## Price Risk

i.e the risk that a VM's price  
will increase relative to others

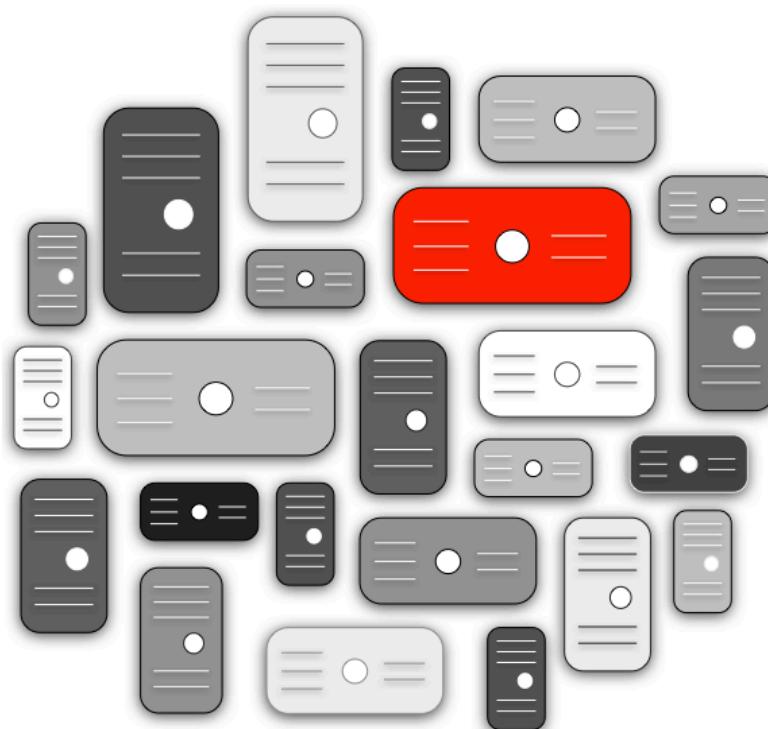
Fault-tolerance  $\cong$  insurance

users pay upfront **premiums** (i.e., fault-tolerance overhead)  
and expect a **payout** later (i.e., ability to limit the loss of work)



*Does mitigating the price  
risk affect **performance**  
and **revocation risk**?*

*How to enable flexible cloud  
applications to mitigate the  
**price risk transparently**?*



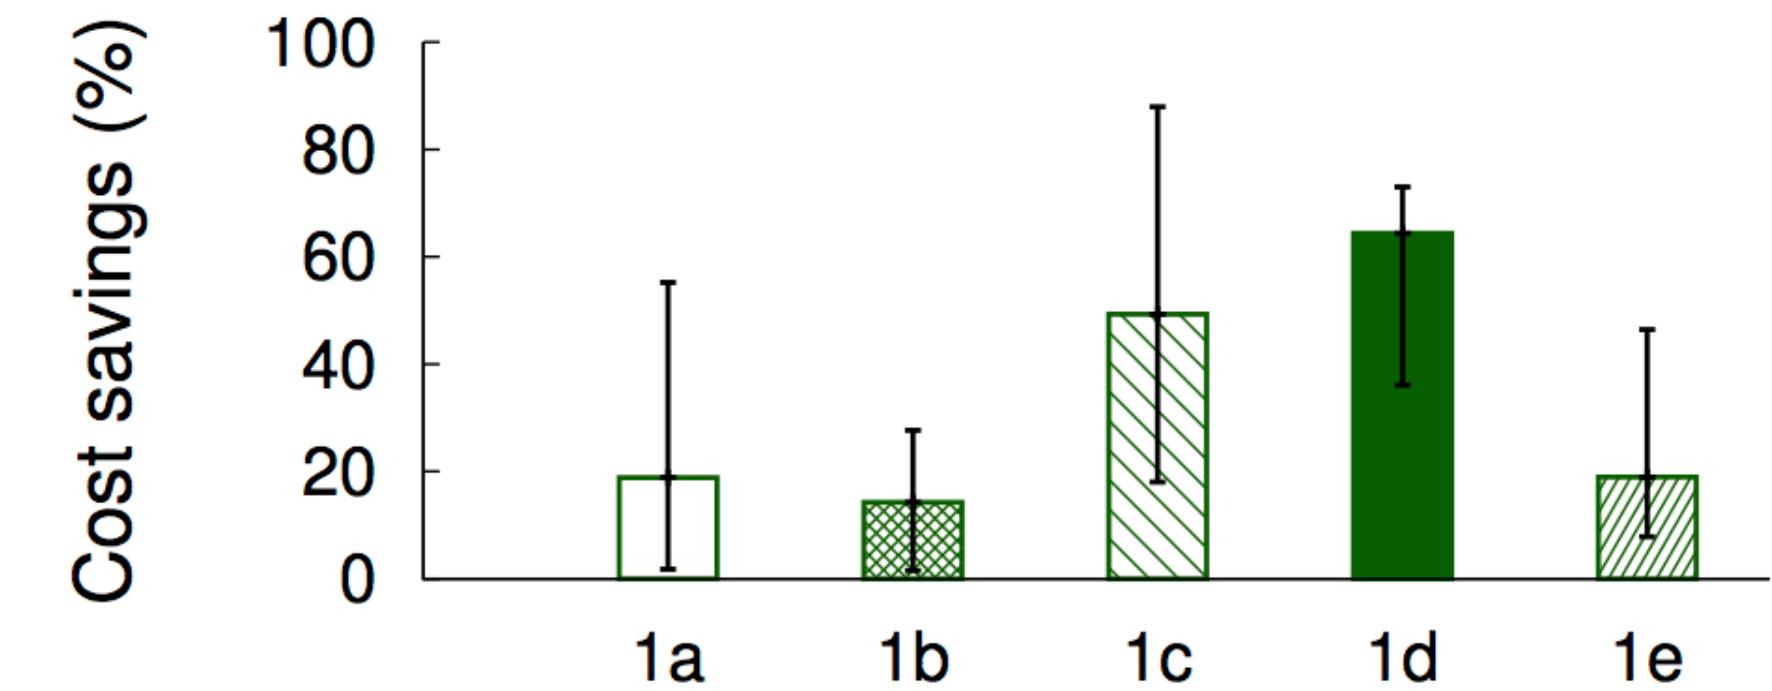
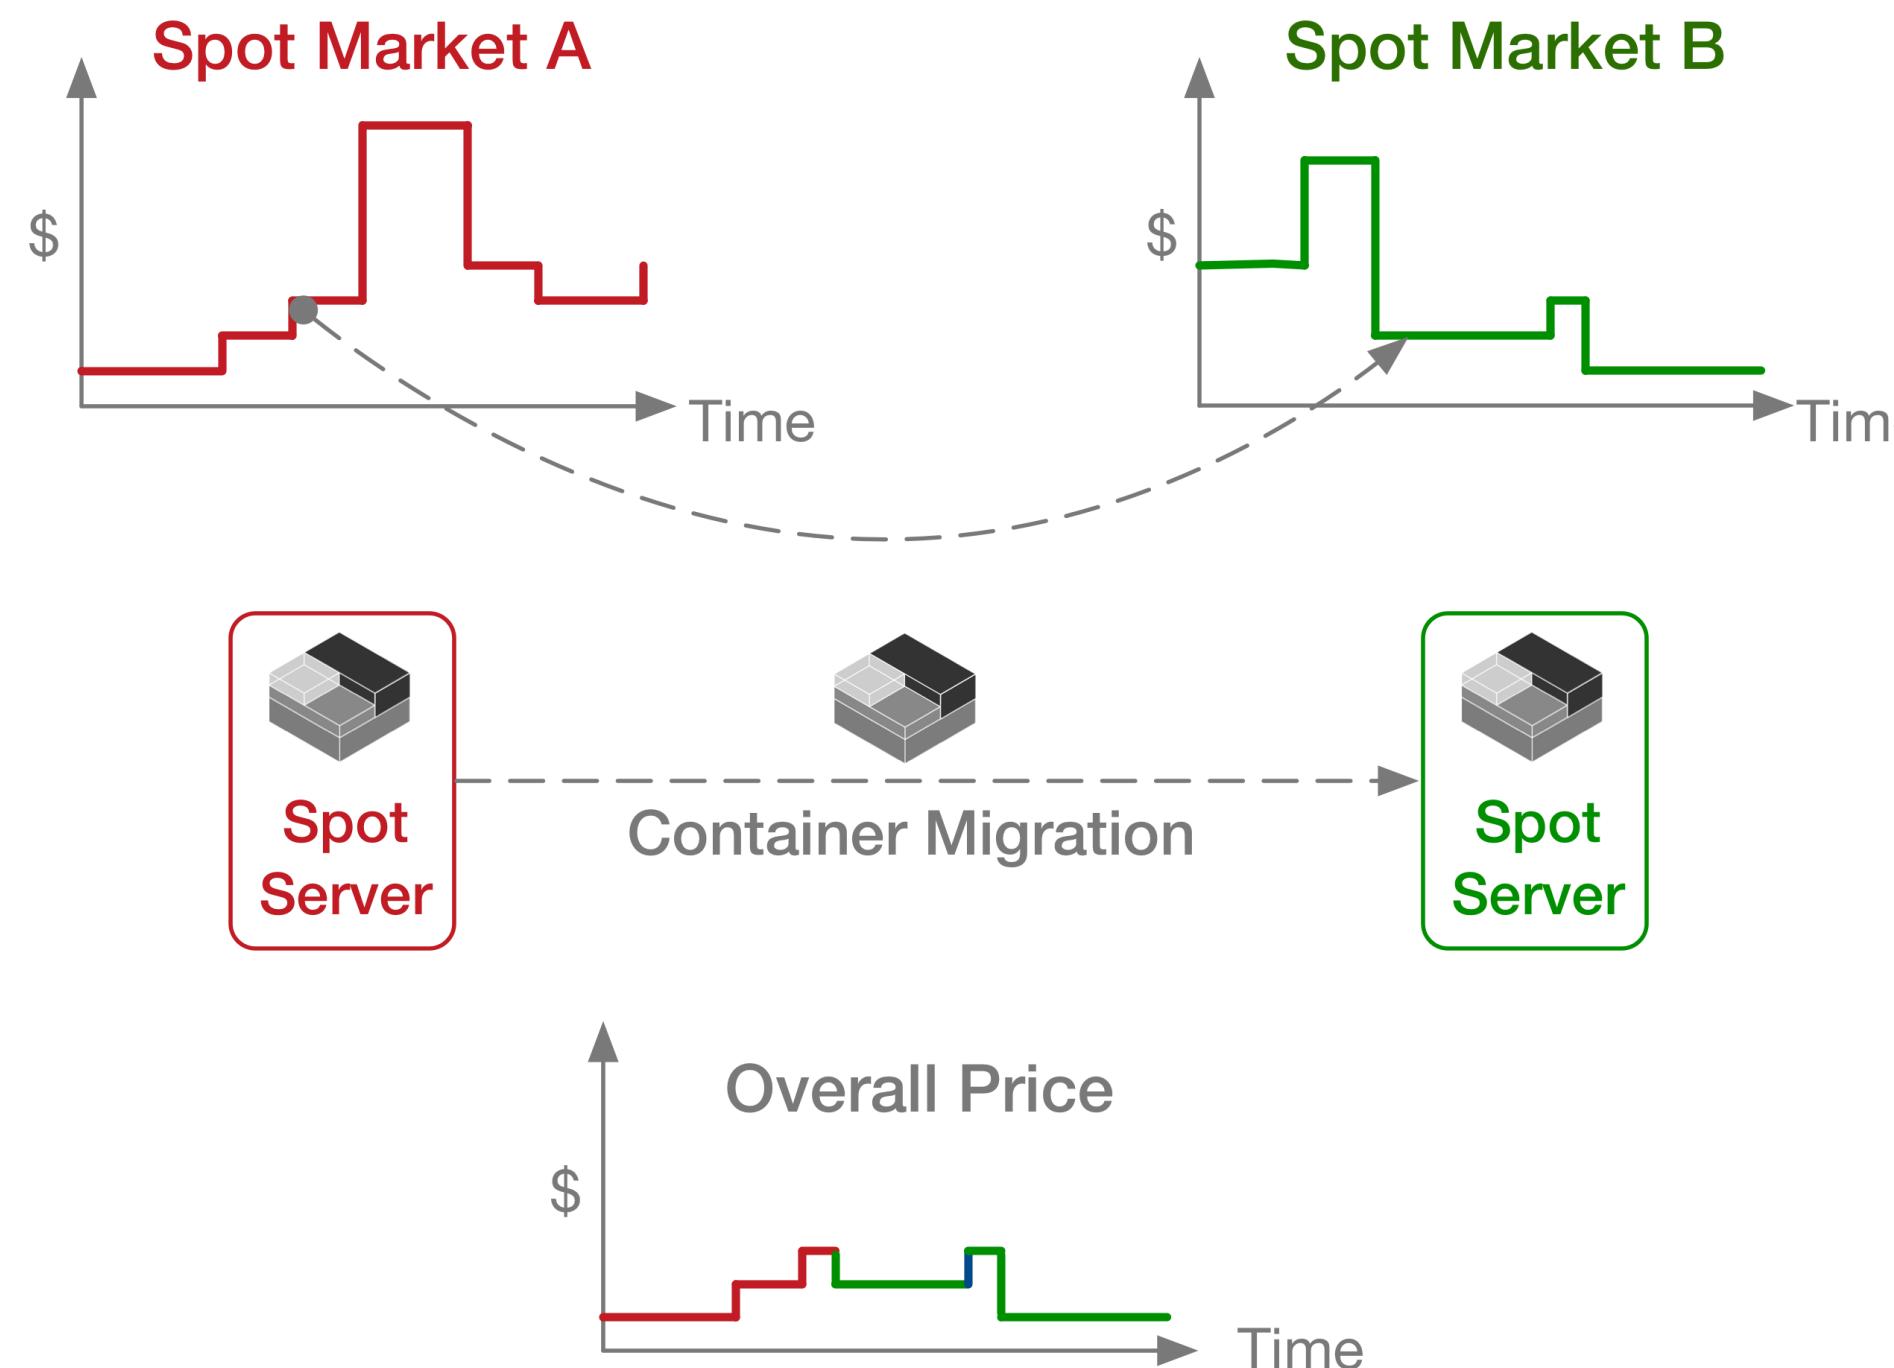
**HotSpot: Automated Server Hopping**

# Automated Server Hopping

“ Change, before you have to ”

A resource container that automatically hops spot VMs as market conditions change

Results from the EC2 spot market  
US-East-1 markets (3/1/2017 - 5/1/2017)

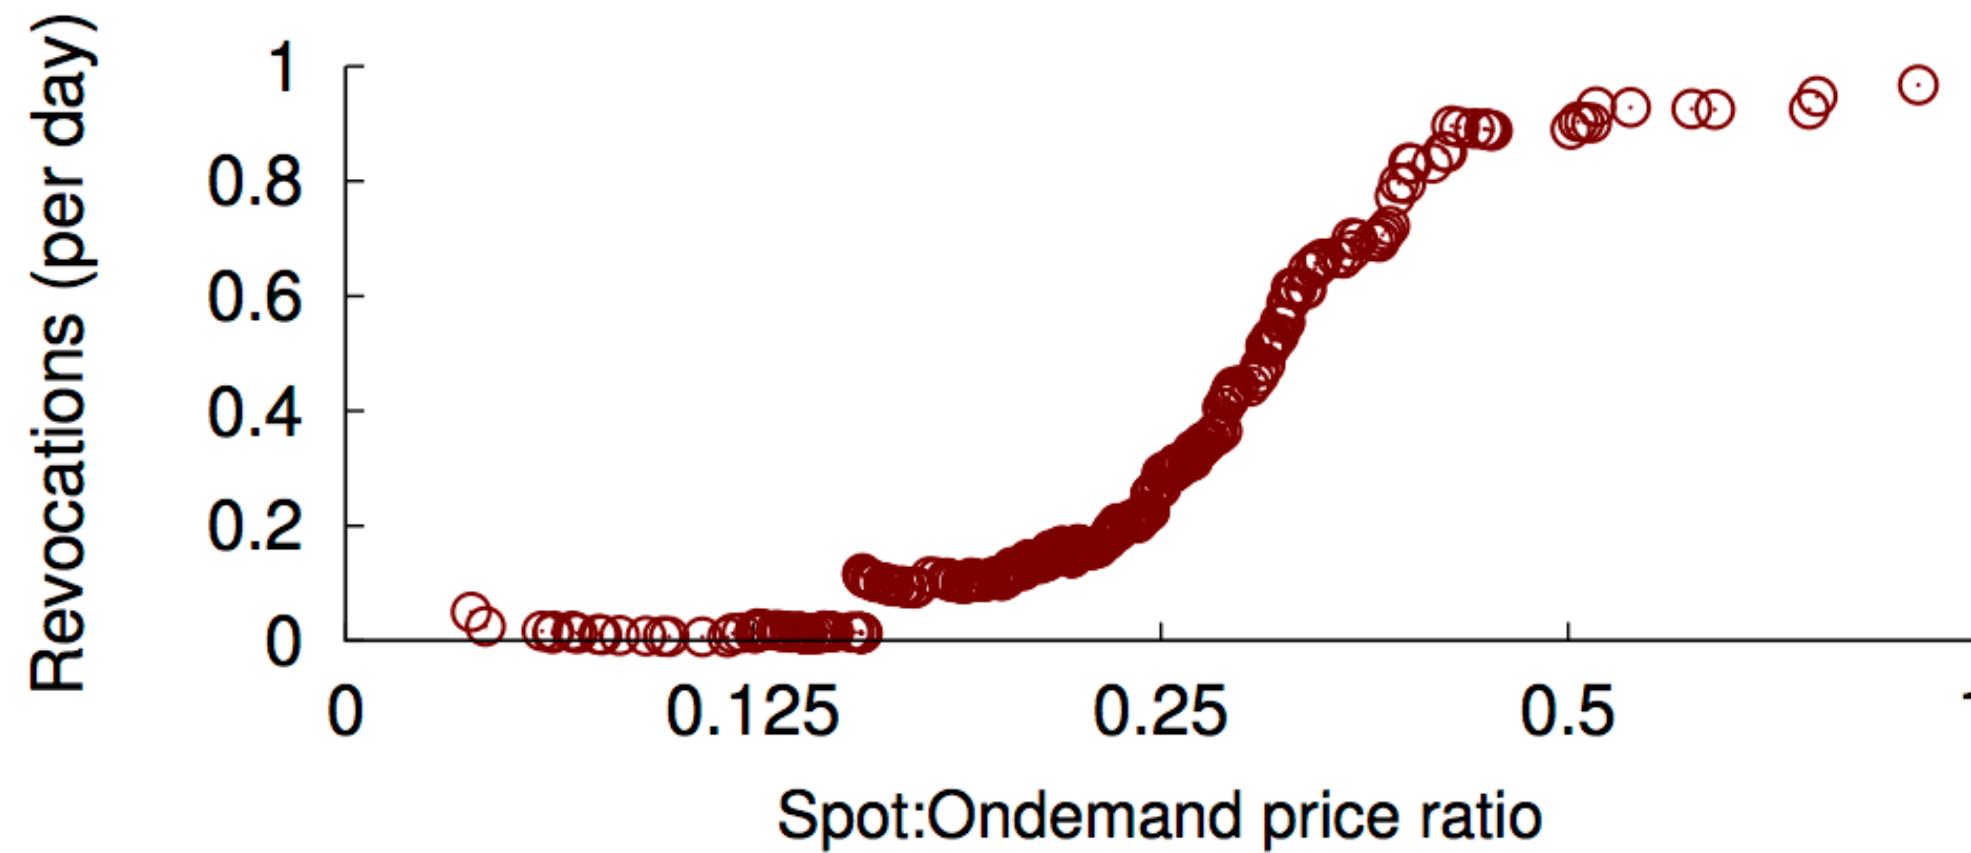


Ideal savings from **hopping** vs. **staying**  
for a long-running job (30 days)

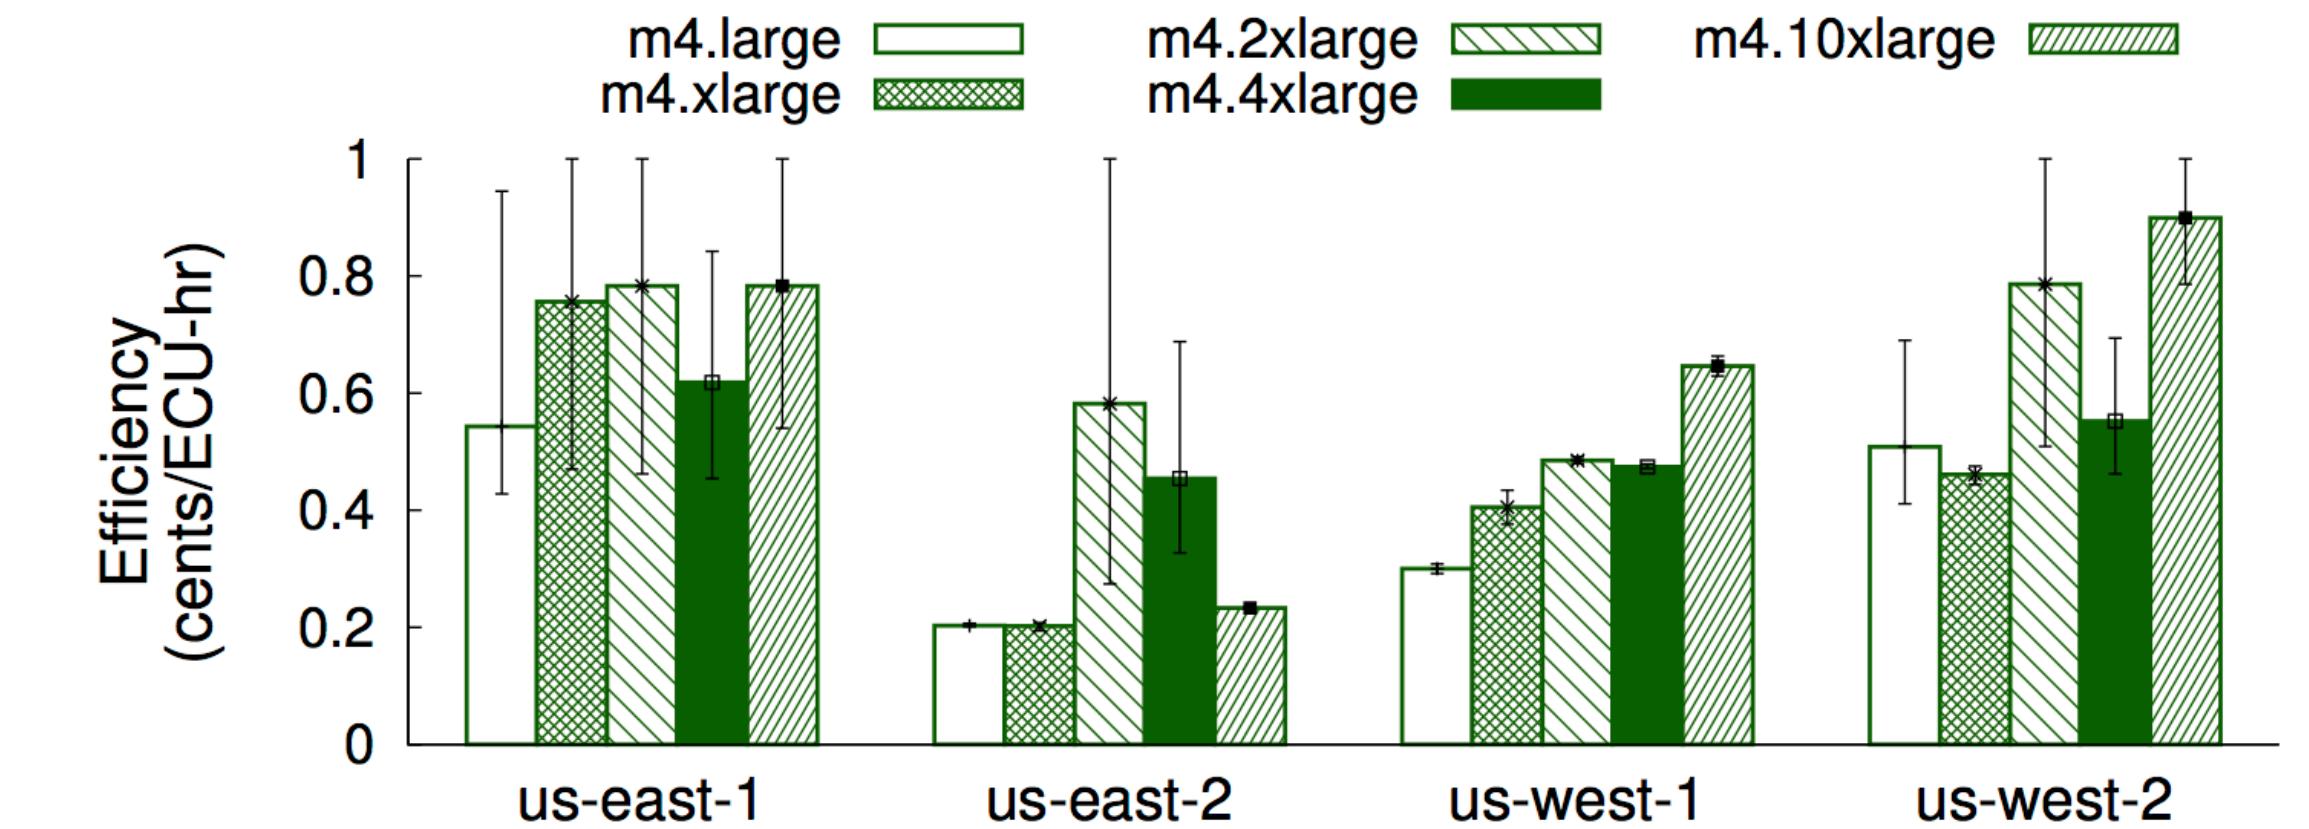
# Effect on Revocation Risk and Performance

Insights from spot market analysis

- 1 Highly discounted servers tend to have lower revocation risk



- 2 Cost efficiency is uncorrelated with VM capacity (and thus performance)



Server hopping lowers revocations without necessarily degrading performance

# Design of Server Hopping Logic

## Migration policy

Run on a VM that has the best cost-efficiency in \$/utilized-resource without hindering the performance

## Policy invariant

**Trigger** a check whenever

- VM utilization changes
- spot market prices change

## Cost-benefit analysis

Migrate to the spot vm that gives the highest cost-benefit gain

### Expected benefit

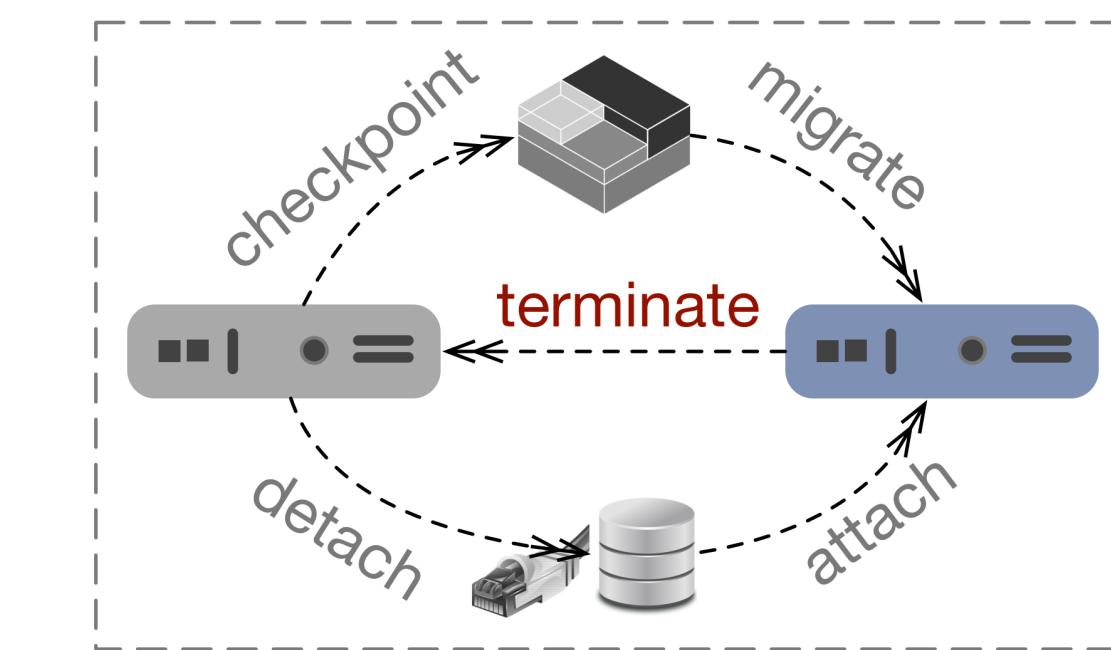
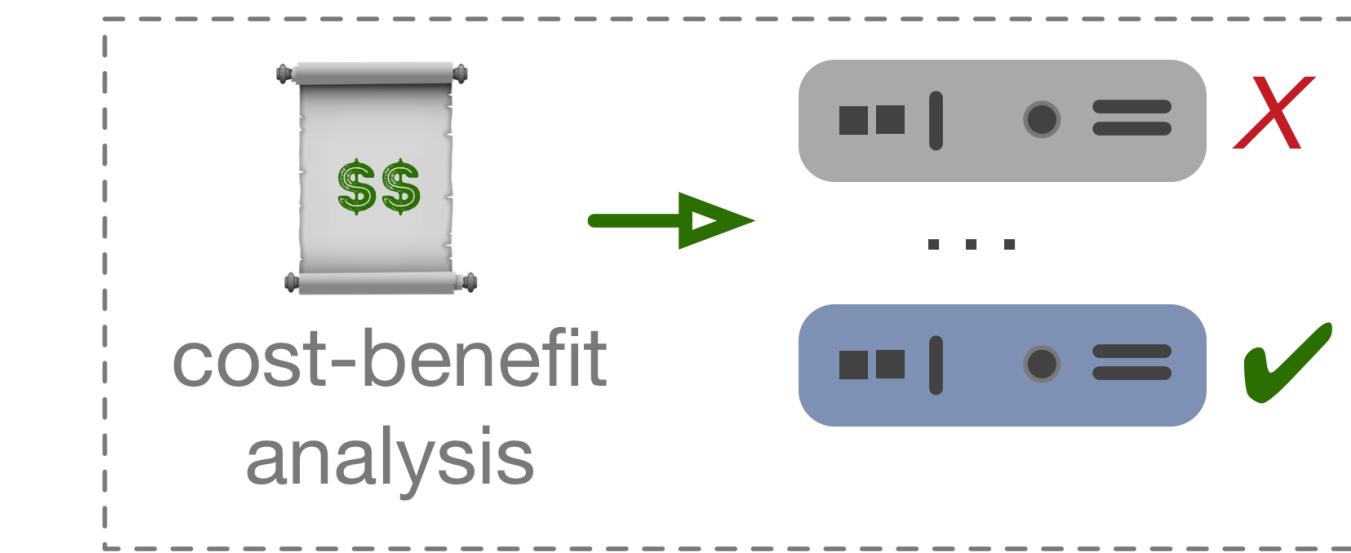
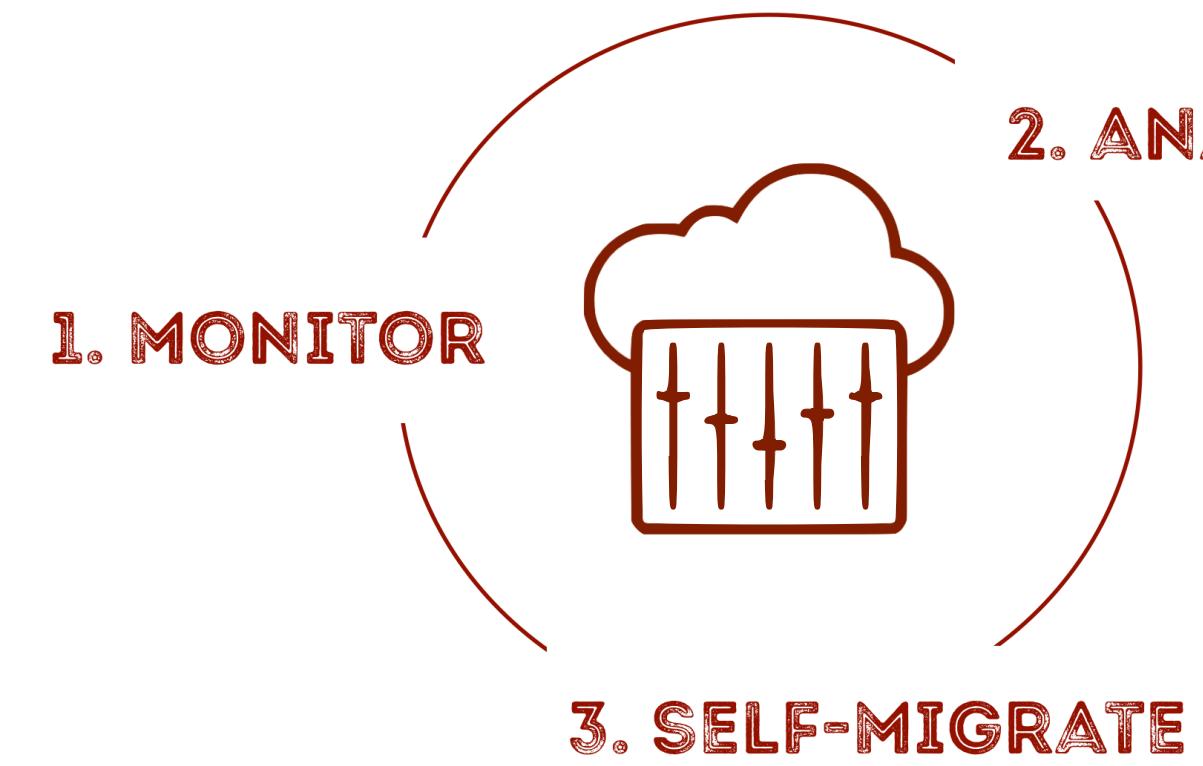
- Gain in cost-efficiency for the duration of expected stay
- $f(\text{market characteristics})$

### Migration cost

- Double-paying for VMs + min. VM holding time
- $f(\text{application footprint})$



# HotSpot: Design and Implementation



Fully functional prototype available at:

<https://sustainablecomputinglab.github.io/hotspot/>

# Evaluation

Compare **cost, performance, revocations** of running a flexible batch application on

**Spot VM**  
**with no protection**  
**(SpotFleet)**

VS.

**Spot VM**  
**with fault-tolerance**  
**(SpotOn [SoCC 2015])**

VS.

**Spot VM**  
**with server hopping**  
**(HotSpot)**

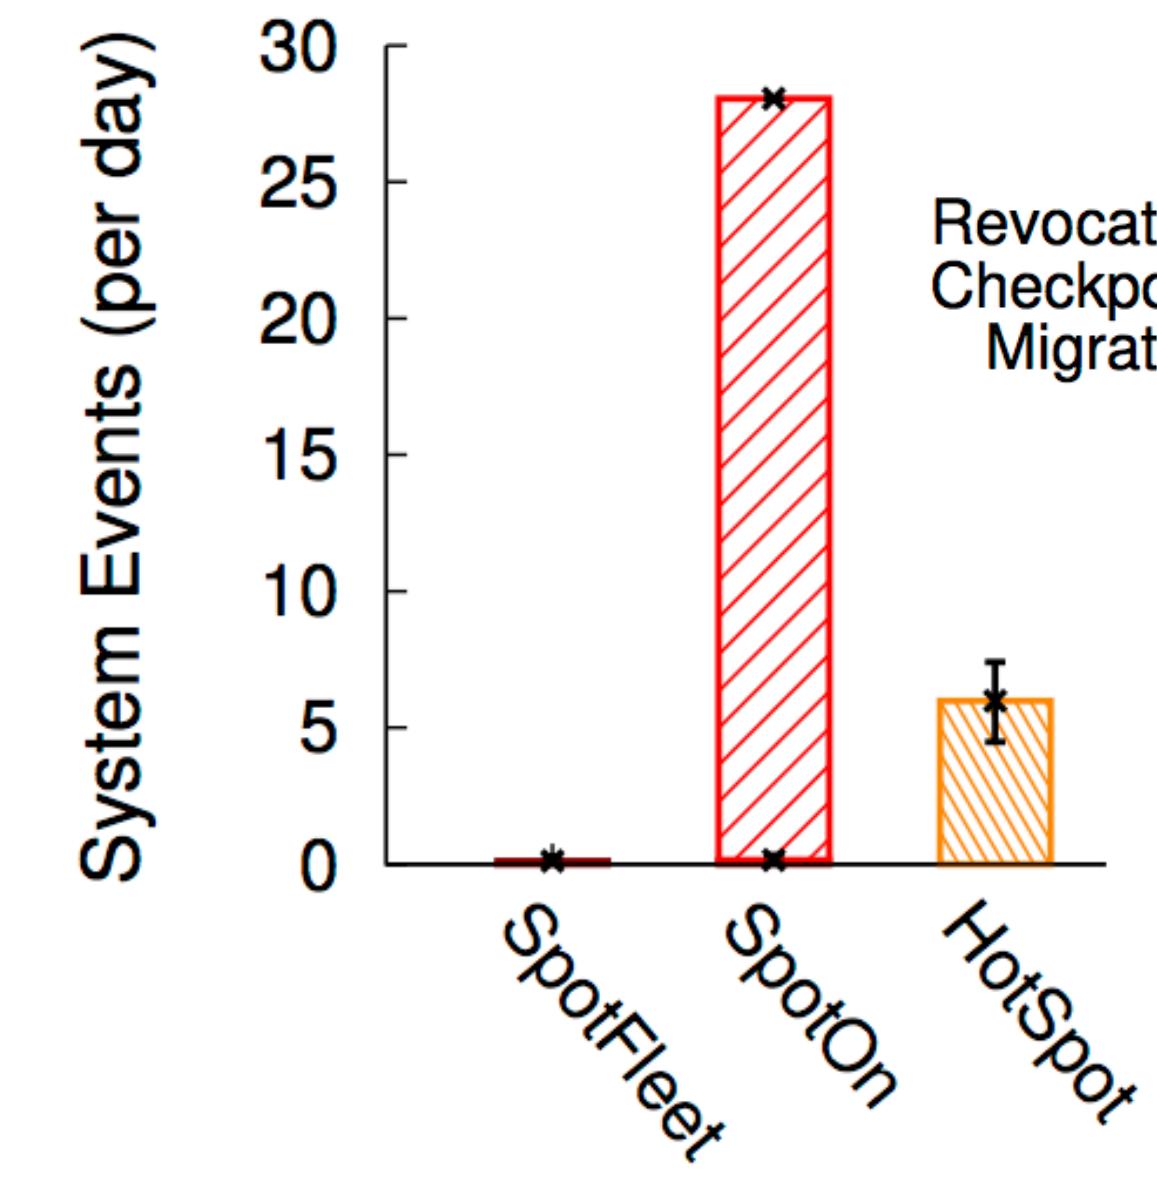
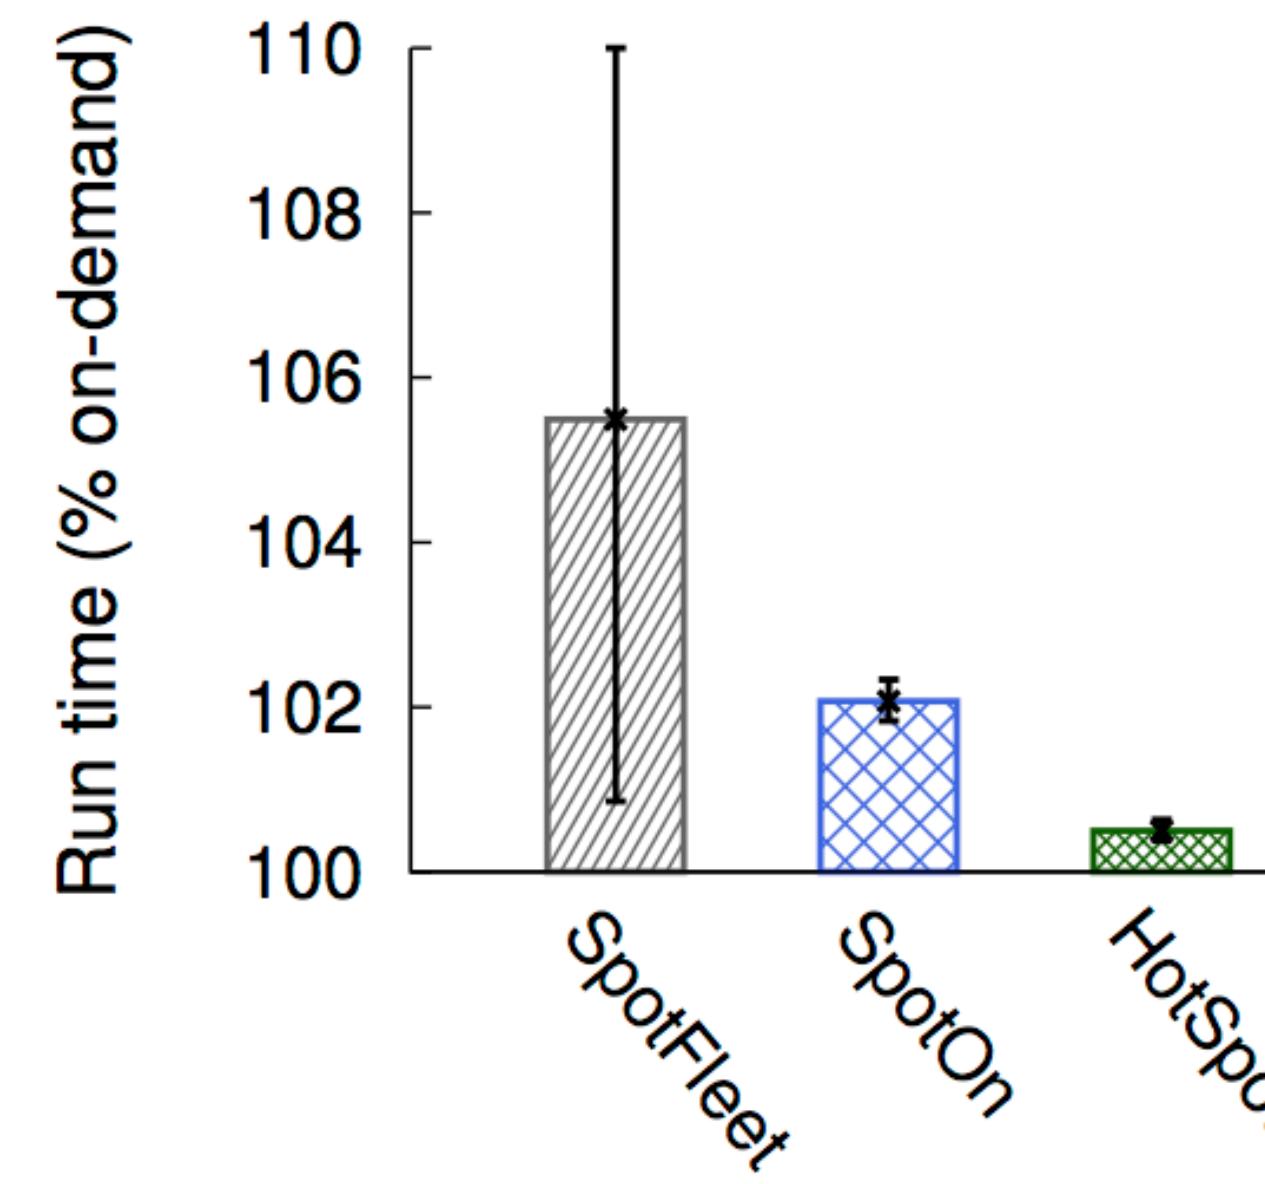
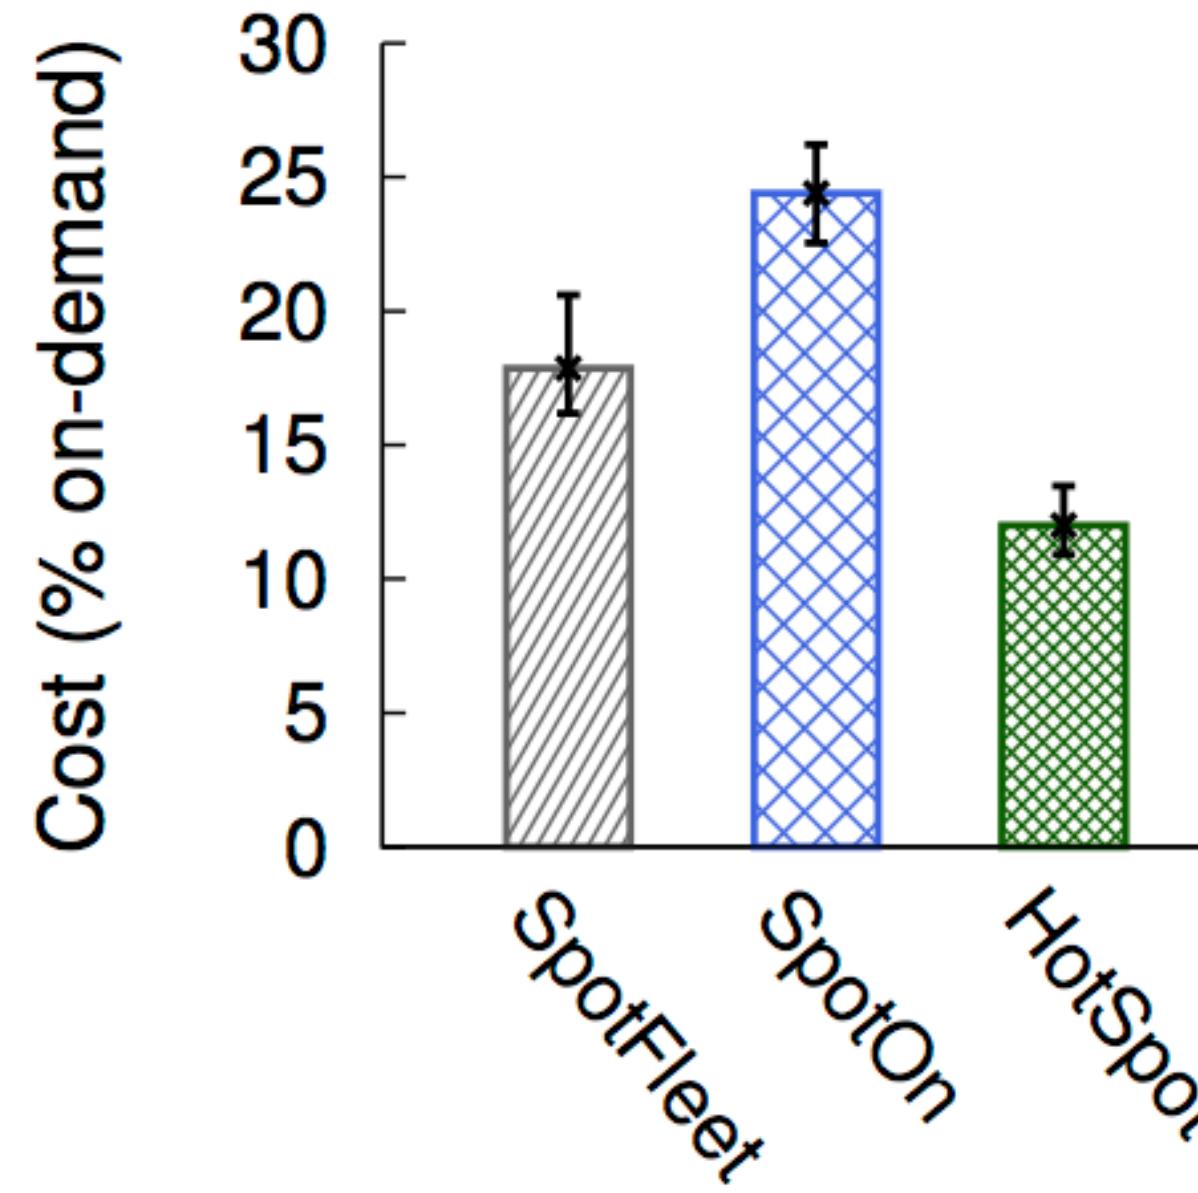
## 1. How do changes in job and market characteristics affect each approach?

Run the prototypes on EC2 (but control job and market conditions using emulators)

## 2. How do different approaches perform on the real market for real jobs?

Simulate running Google cluster trace jobs on Amazon spot price traces (03/2017 to 05/2017)

# Google Cluster Traces on EC2 Spot Markets



Even in the current EC2 spot markets (with low revocation rates),  
optimizing for price-risk results in **30-50% additional savings** without degrading performance

# Conclusion

Transient server markets are an emerging area and offer many opportunities for cost savings

## Price Risk



Price risk is significant in current spot markets

Mitigating price risk also reduces revocations

## HotSpot



Proposed the technique of **automated server hopping**

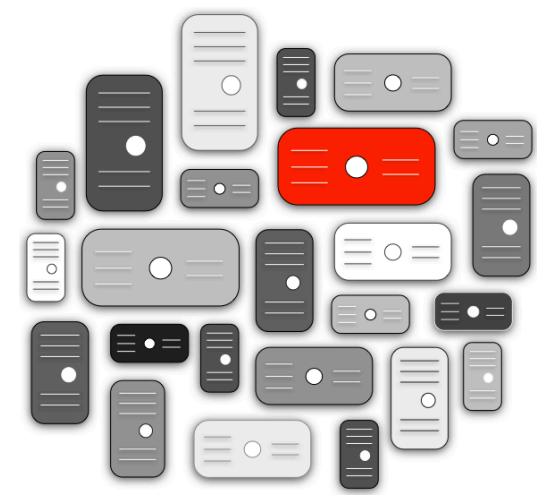
Designed and implemented HotSpot for EC2 spot markets

## Evaluations

30-50% COST REDUCTION

vs. other techniques

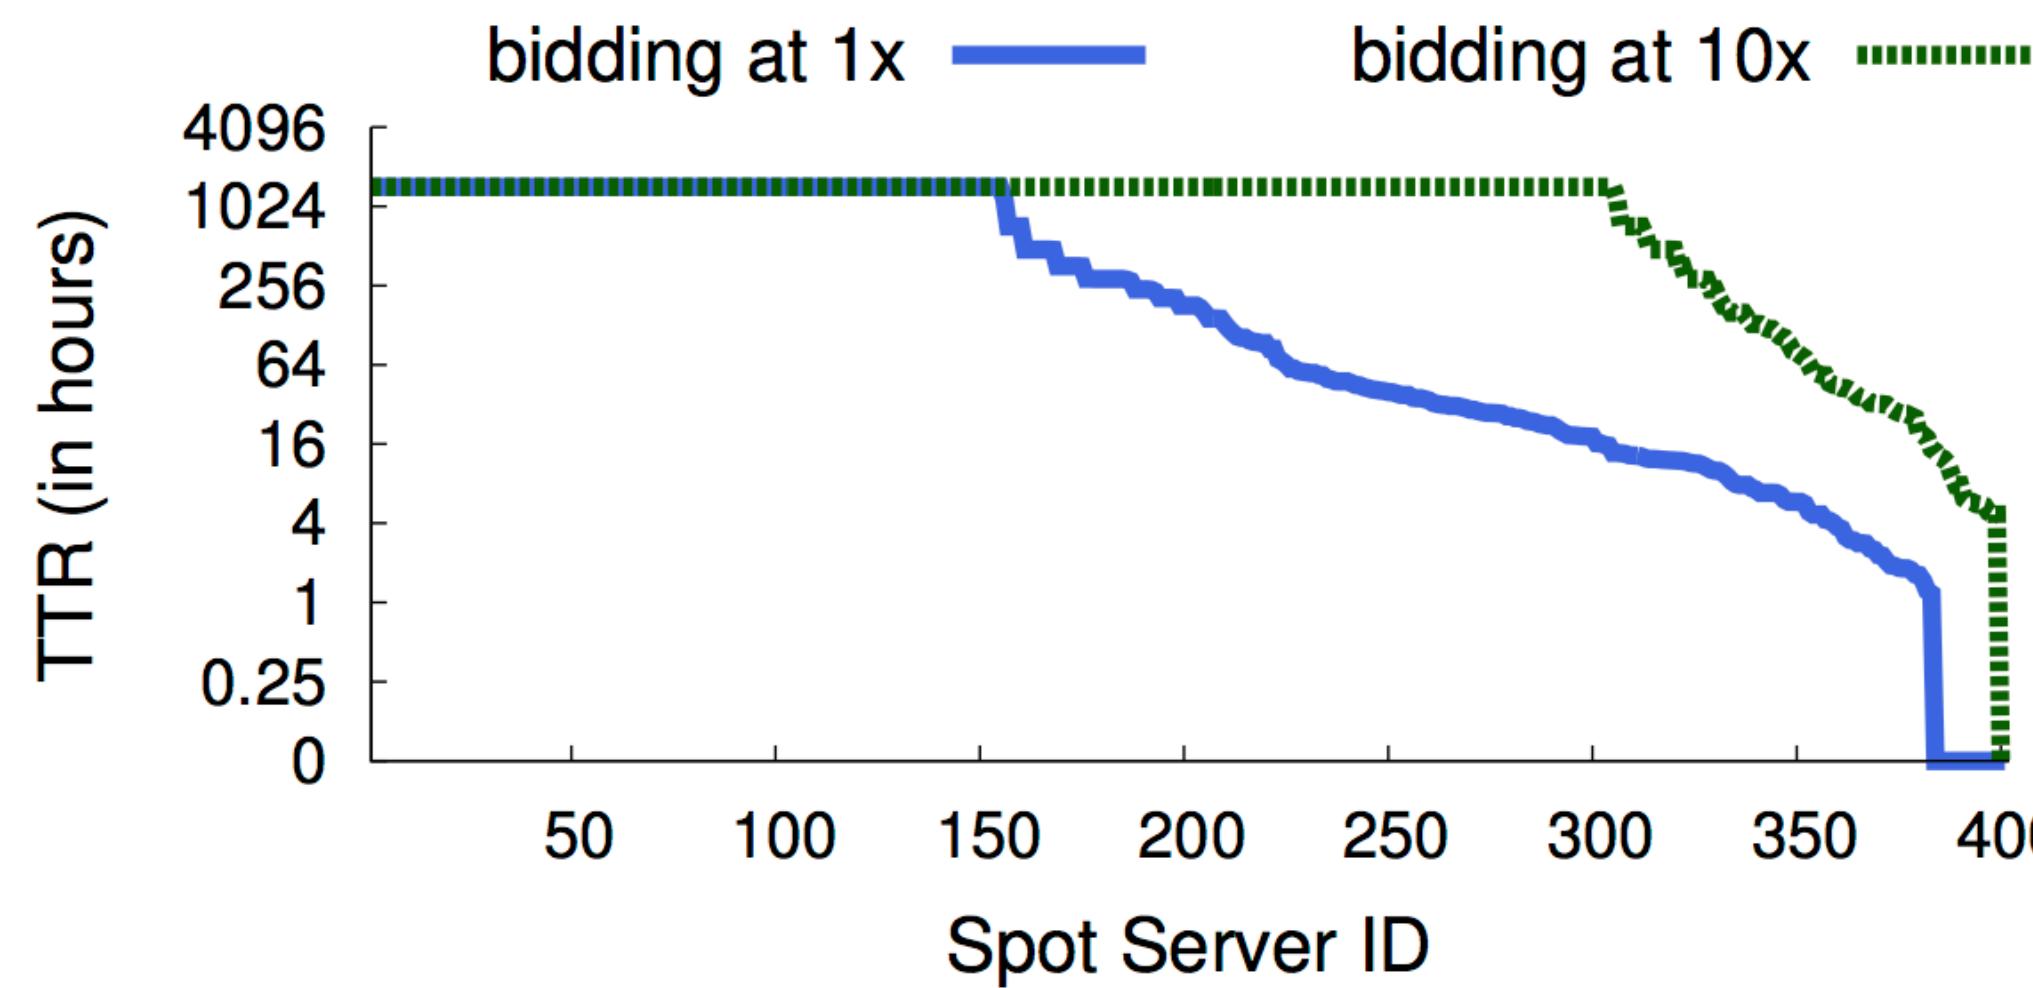
- Lower Overhead
- Lower Revocations
- More Deterministic



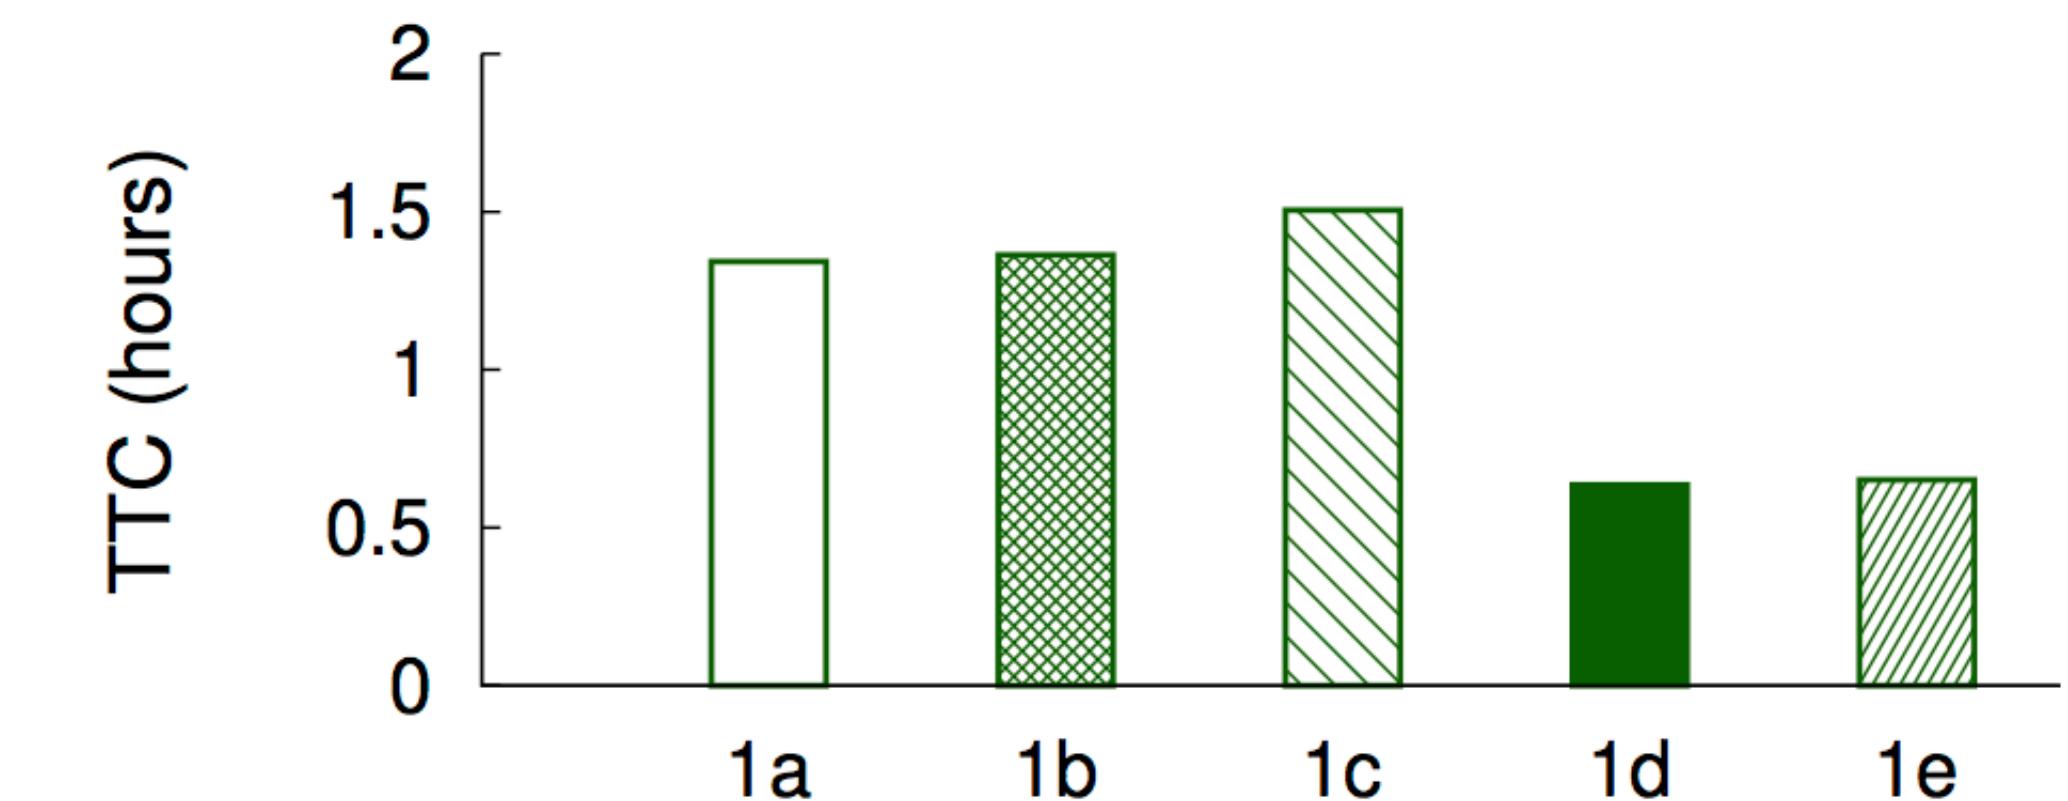
# Backup Slides

# Price Risk >> Revocation Risk

Data from all 402 spot VMs in US-East-1 over 3/1/2017 to 5/1/2017



Mean Time-to-Revocation (TTR) when  
bidding 1x is ~25 days and 10x is ~47 days

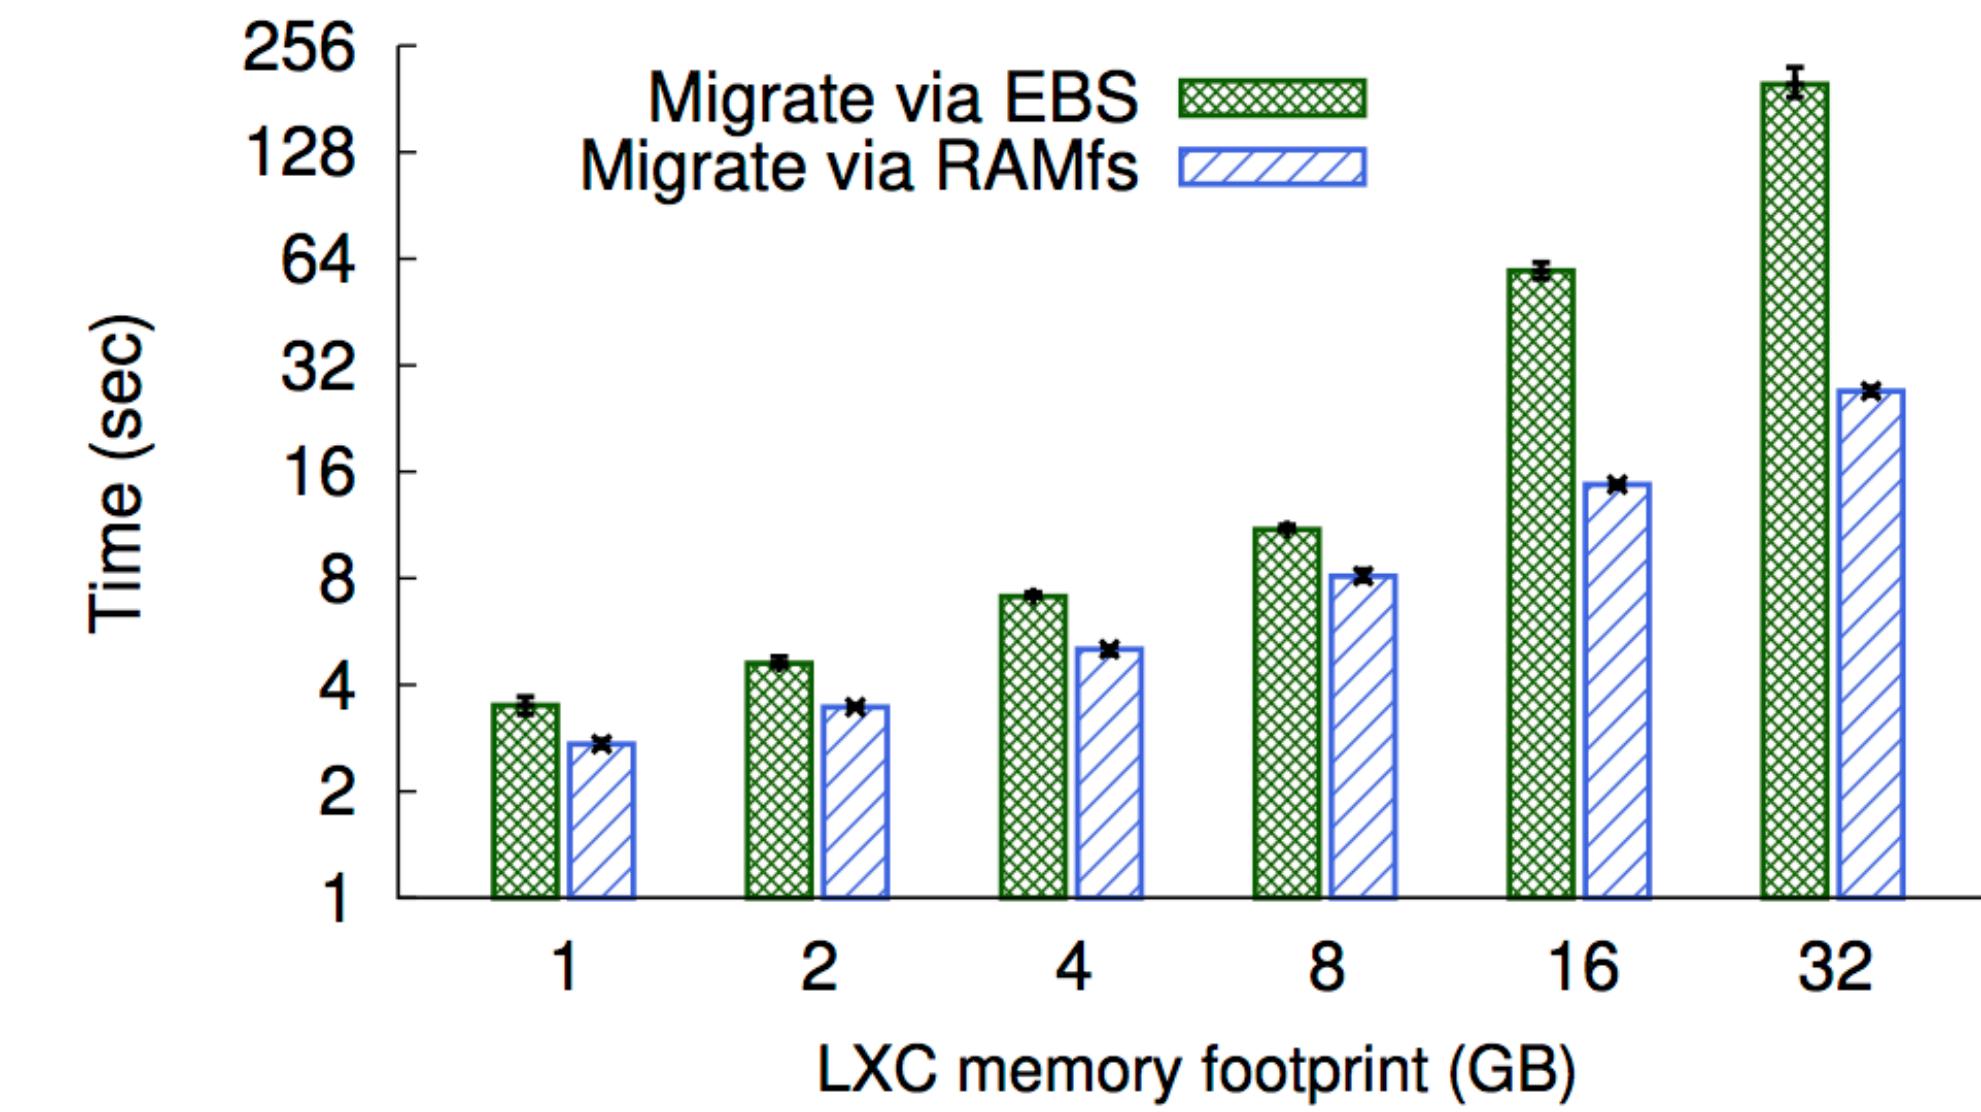


Time-to-Change (TTC) for the  
cheapest VM is **1.1 hours**

# Migration Latencies in EC2

Operation	Min (sec)	Mean (sec)	Max (sec)
<i>Price and Resource Monitoring</i>	<1	<1	<1
<i>Acquire On-demand VM</i>	16	28	31
<i>Acquire Spot VM</i>	31	67	167
<i>Transferring Disk &amp; Network</i>	18	28	48
<i>Terminate Source VM</i>	31	44	46
<b>Total</b>	~64-80	~101-140	~126-262

Platform's API operations

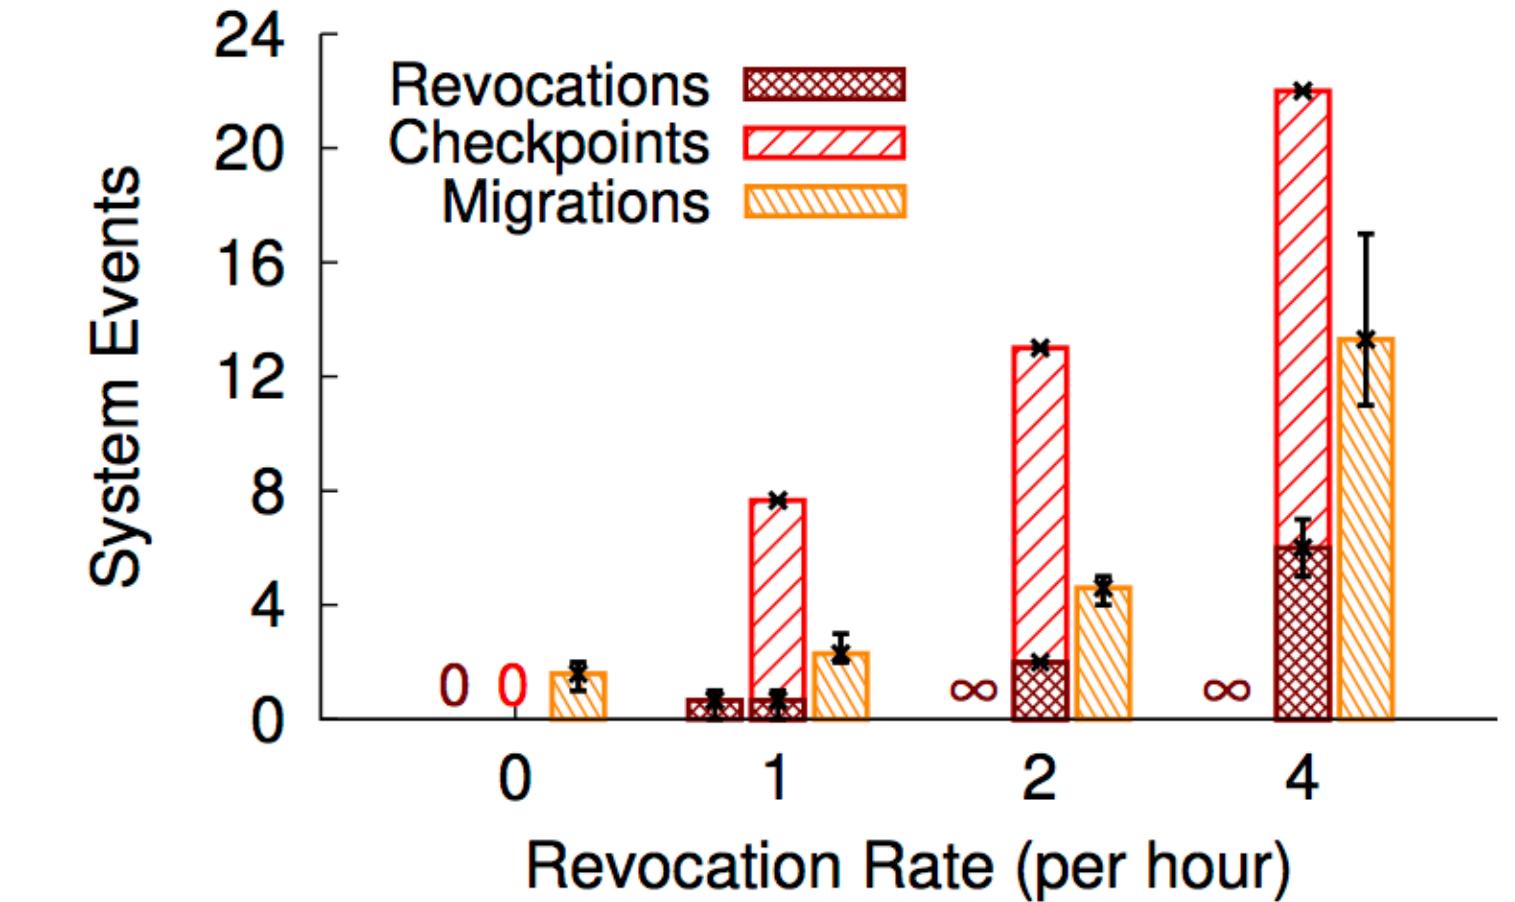
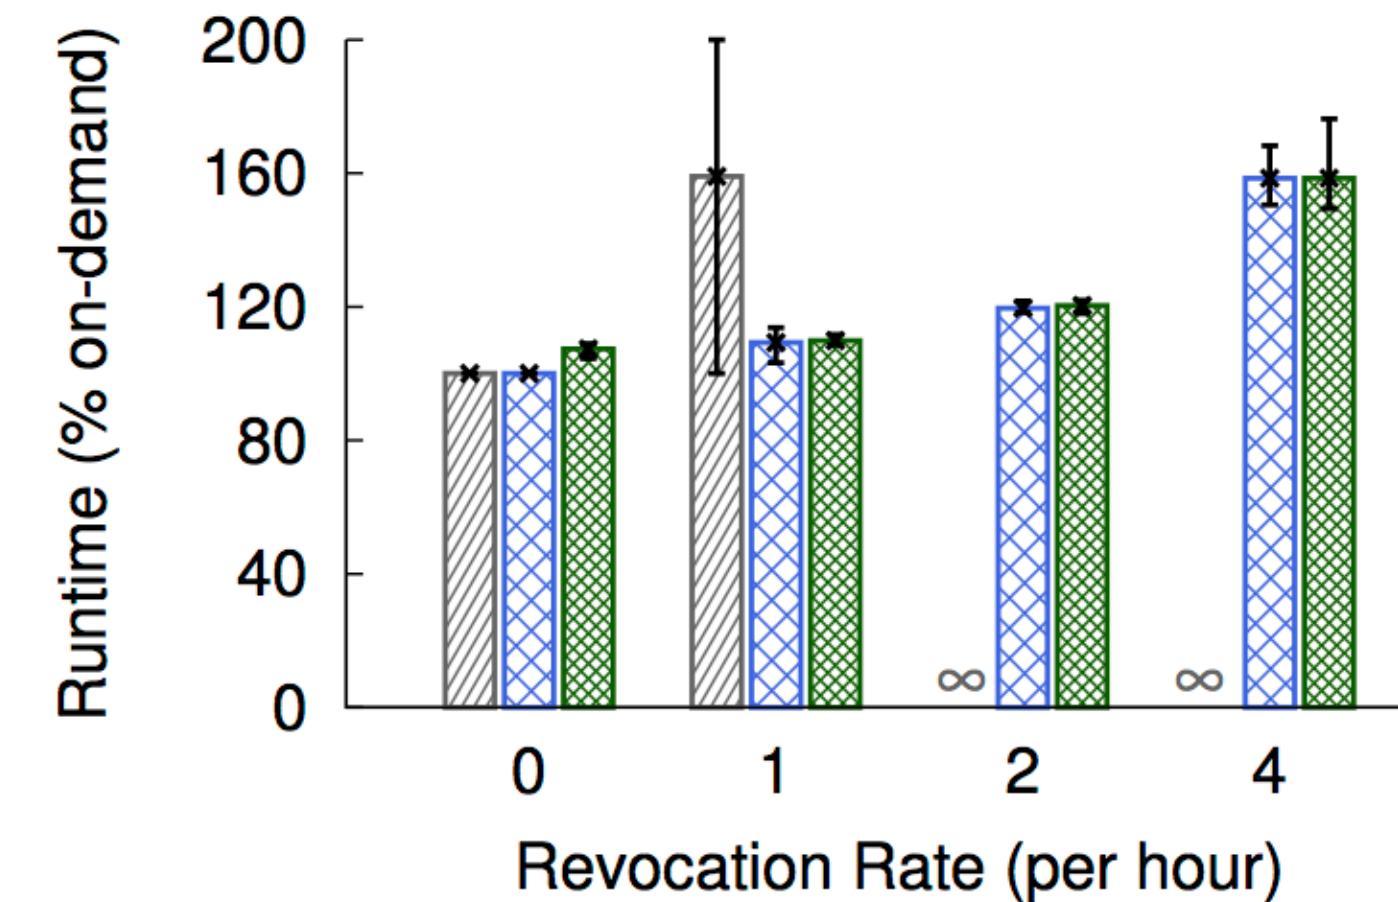
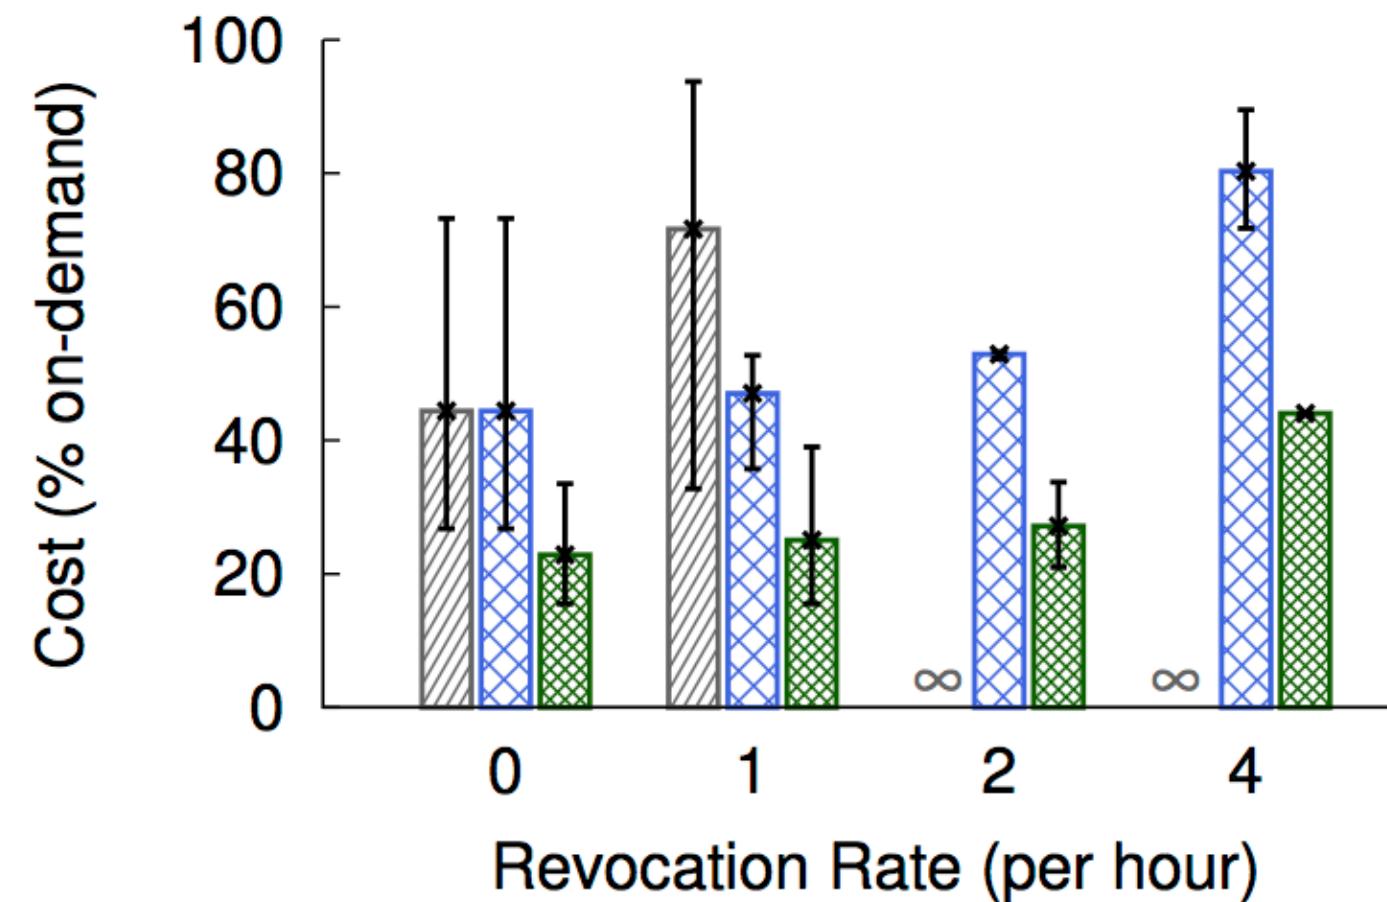


# Effect of Changes in Market Volatility

SpotFleet

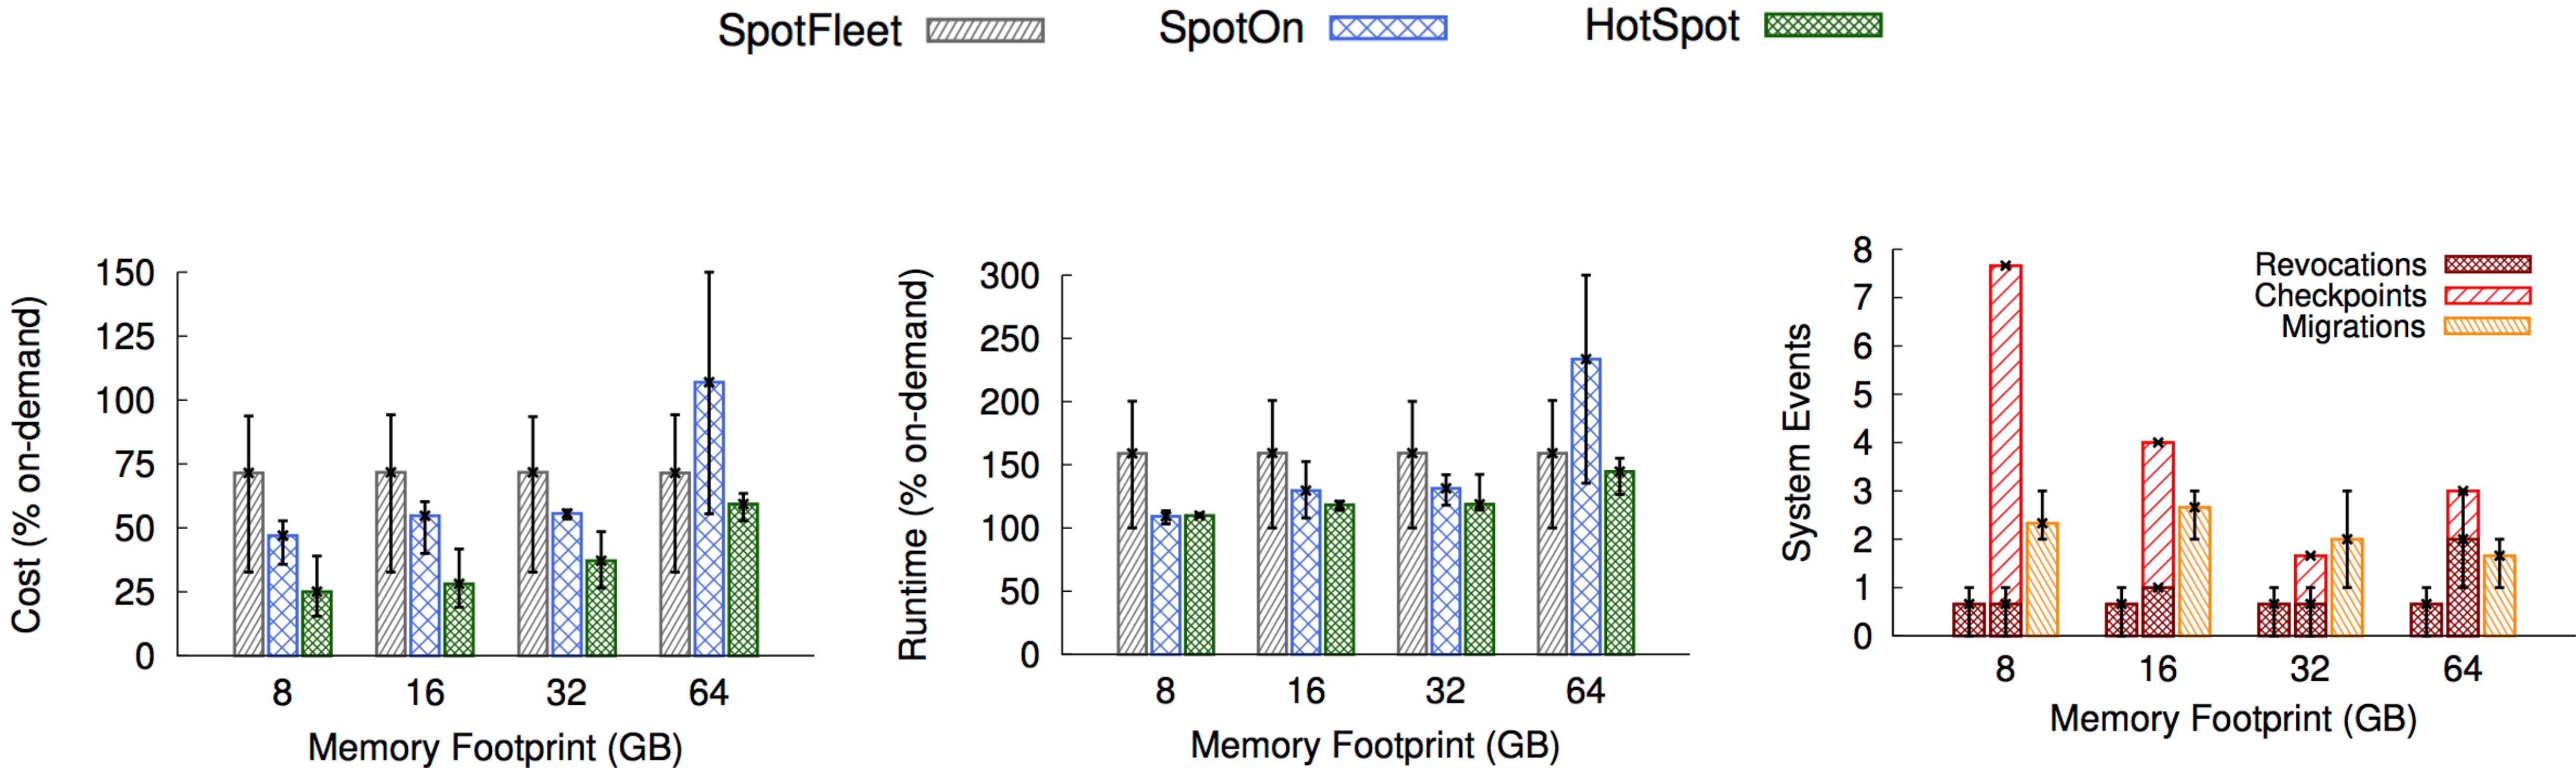
SpotOn

HotSpot



As markets become more volatile,  
HotSpot's savings will improve relative to SpotFleet and SpotOn

# Effect of Changes in App Footprint



HotSpot outperforms both SpotFleet and SpotOn at all levels, though its gains reduce with increase in the memory footprint.