

CS3640

Transport Layer (4): TCP

Prof. Supreeth Shastri

Computer Science

The University of Iowa

Lecture goals

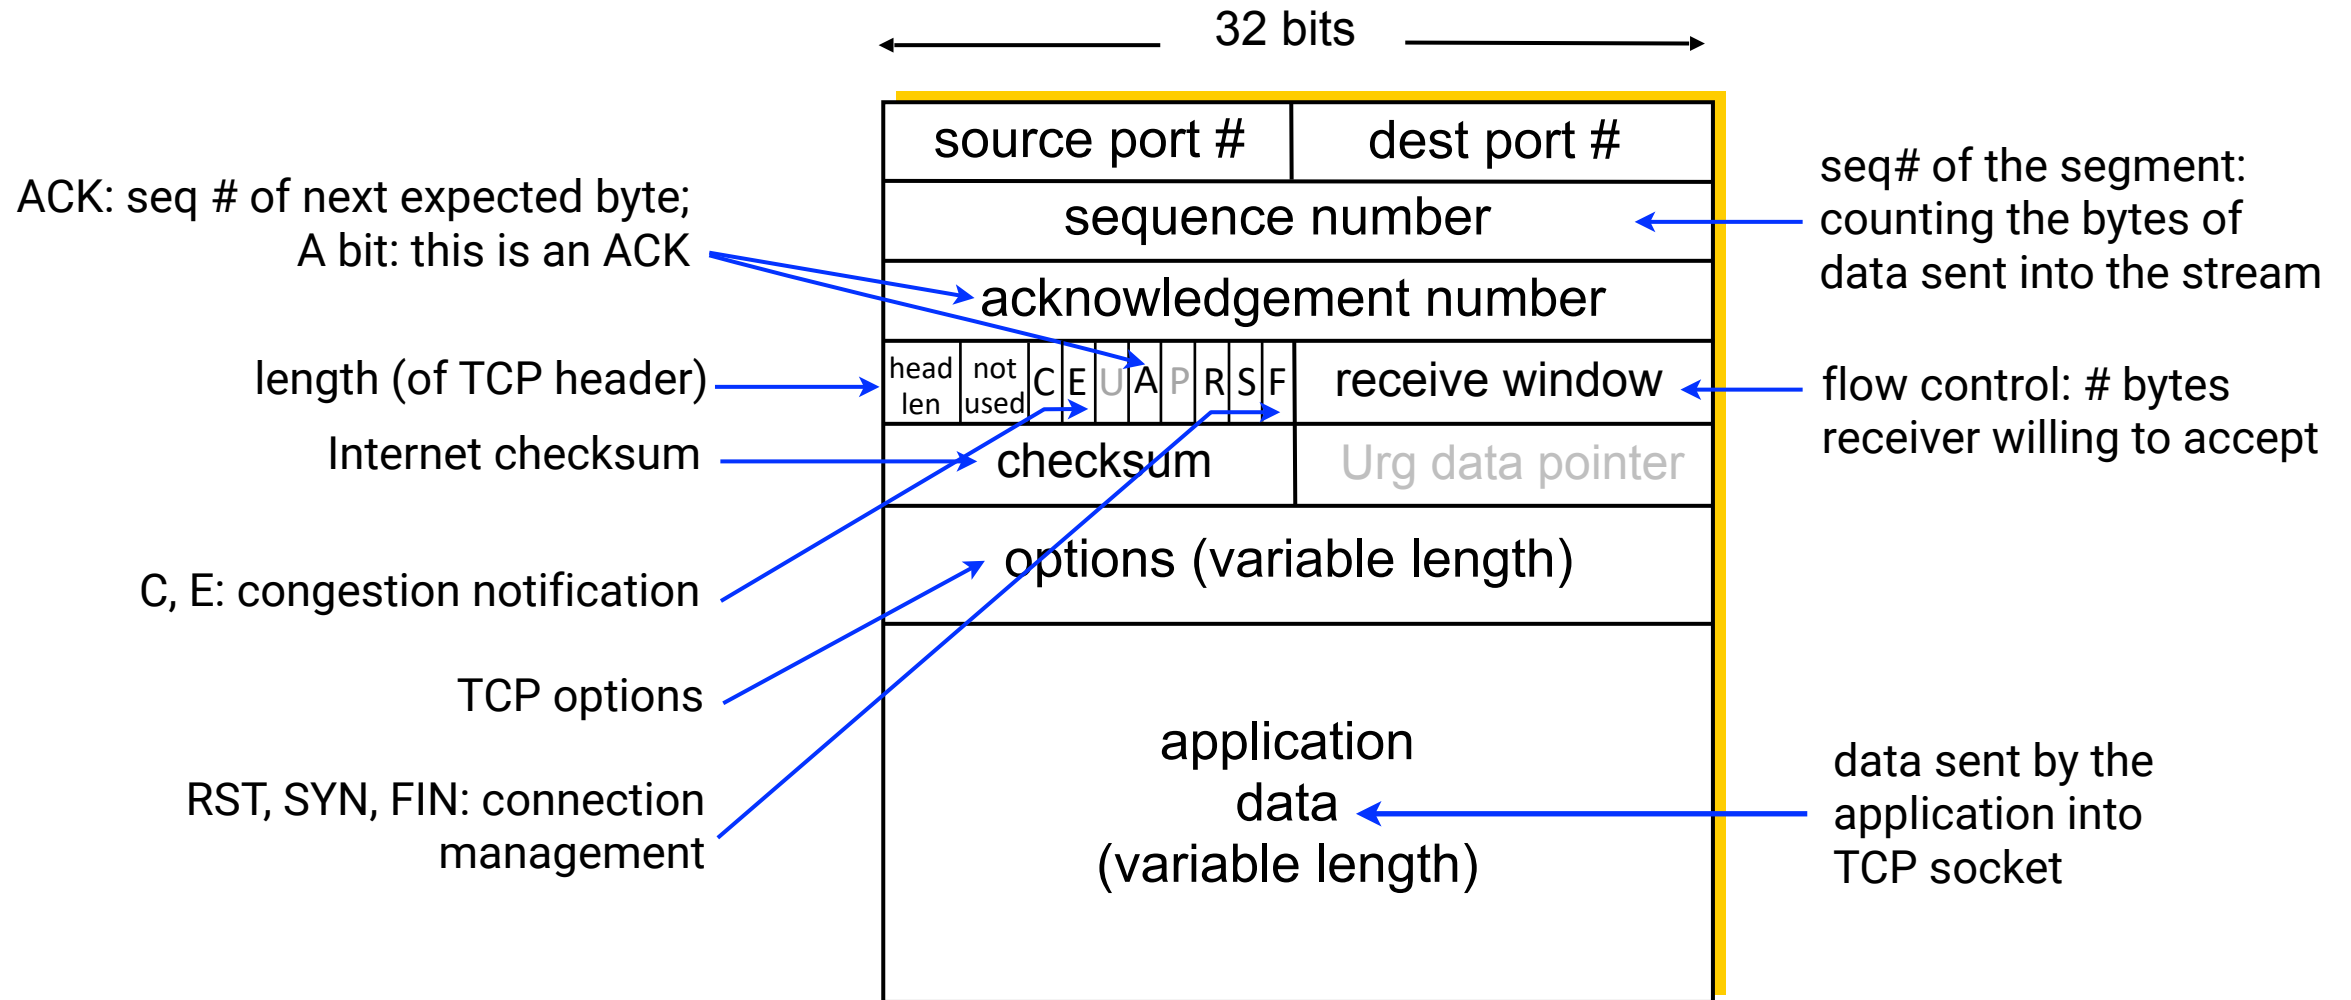
from principles to practice: design and operation of TCP

- *Protocol structure*
- *Connection management*
- *Reliable data transfer*
- *Flow and congestion control*



Chapters 3.5, 3.7

Structure of the TCP segment

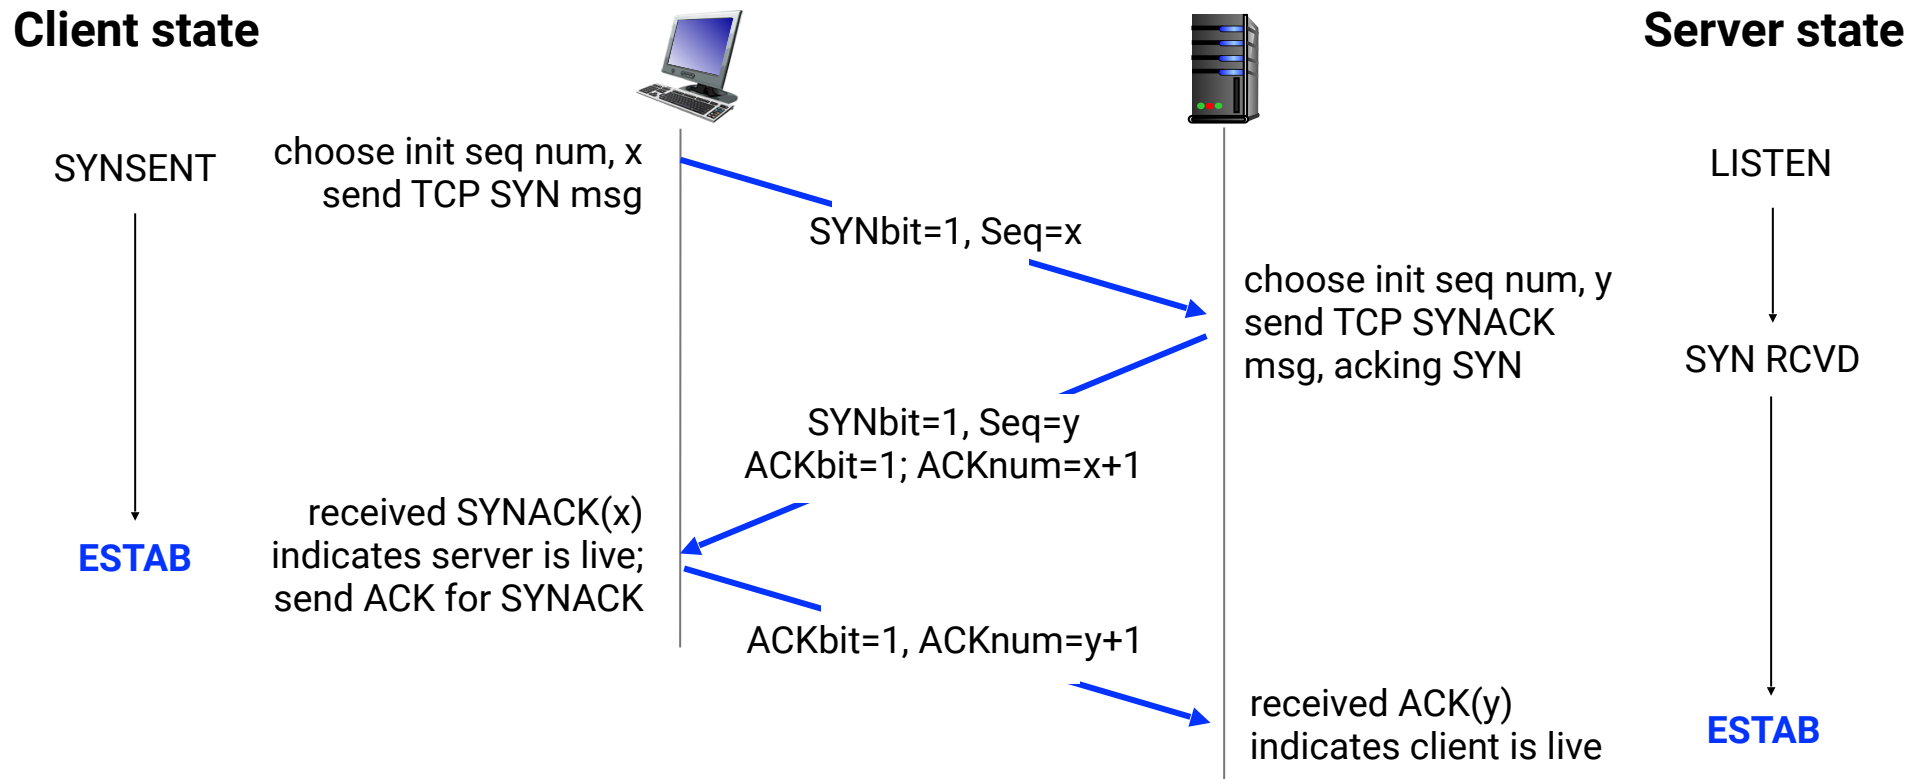


TCP Connection Management

TCP 3-way handshake

TCP is connection-oriented, thus needs a “handshake” before exchanging data

- Goal-1: sender and receiver determine that the other side is willing to establish connection
- Goal-2: sender and receiver agree on connection parameters (e.g., starting sequence #)

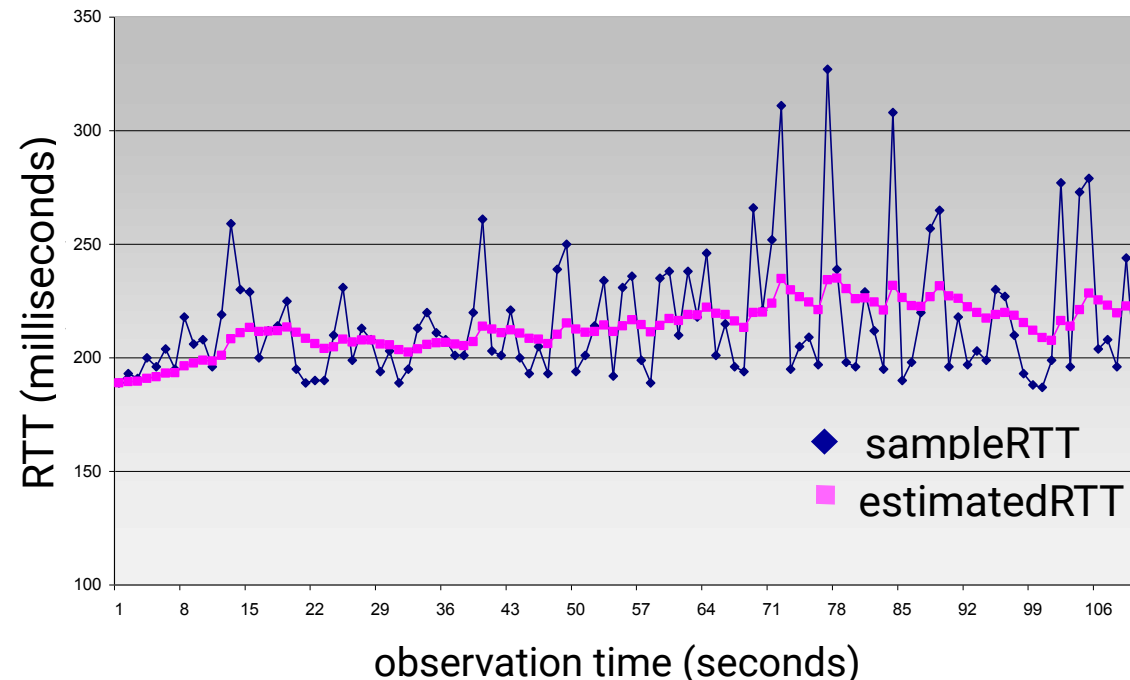


Round Trip Time (RTT) and TCP Timeout

- Key idea: measure time between segment transmission until its ACK receipt
- Such **SampleRTT** will vary over time, so we want estimated RTT to be “smoother”

$$\text{EstimatedRTT} = (1-\alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- exponential weighted moving average (EWMA)
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$



Round Trip Time (RTT) and TCP Timeout

- Underestimating timeout value \Rightarrow unnecessary retransmissions;
Overestimating timeout value \Rightarrow slower loss recovery
- TCP computes timeout interval using **EstimatedRTT**
- Larger the variation in **EstimatedRTT**, larger the safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



safety margin

- **DevRTT**: EWMA of **SampleRTT** deviation from **EstimatedRTT**

$$\text{DevRTT} = (1 - \beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

TCP Reliable Transfer

TCP is a hybrid between GBN and SR protocols

	Go-Back-N	Selective Repeat
ACKs	Cumulative i.e., ACK(k) will ACK all packets up to and including #k	Individual i.e., ACK(k) just ACKs packet #k
Out of order packets	Receiver discards all out of order packets	Buffers out-of-order packet for later delivery
Buffer size	Sender buffer = N; receiver buffer = 1	Sender buffer = N; Receiver buffer = M
Sender timer	Set for only the oldest unacknowledged packet	Set for every transmitted packet

Understanding Sequence and ACK Numbers

Sequence numbers


byte stream “number” of first byte
in segment’s data

Acknowledgement numbers

sequence# of next byte expected
from the other side

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer




sent ACKed sent, not-yet ACKed ("in-flight") usable but not yet sent not usable

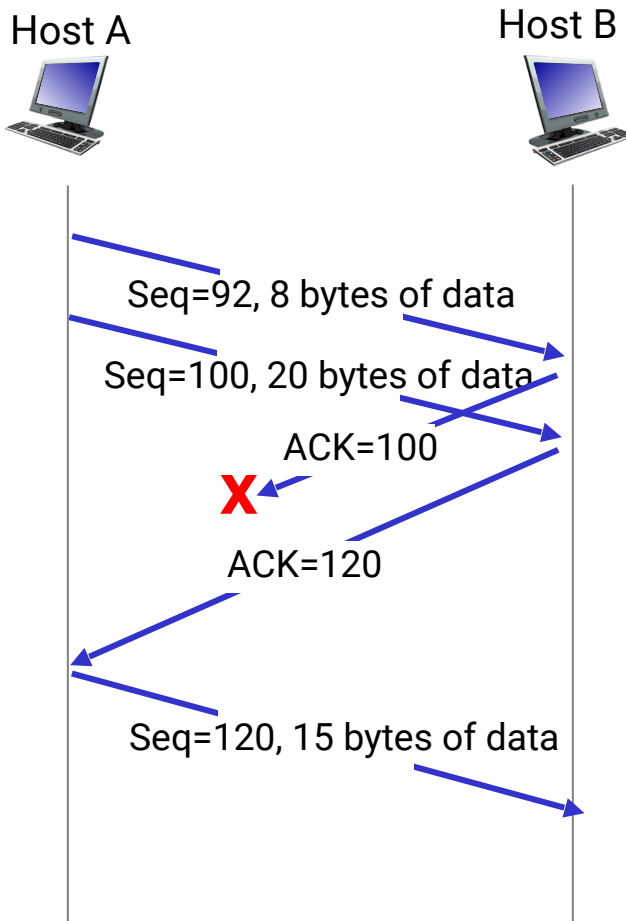
outgoing segment from receiver

source port #	dest port #
sequence number	
	A rwnd
checksum	urg pointer


Example Retransmission Scenarios



lost ACK

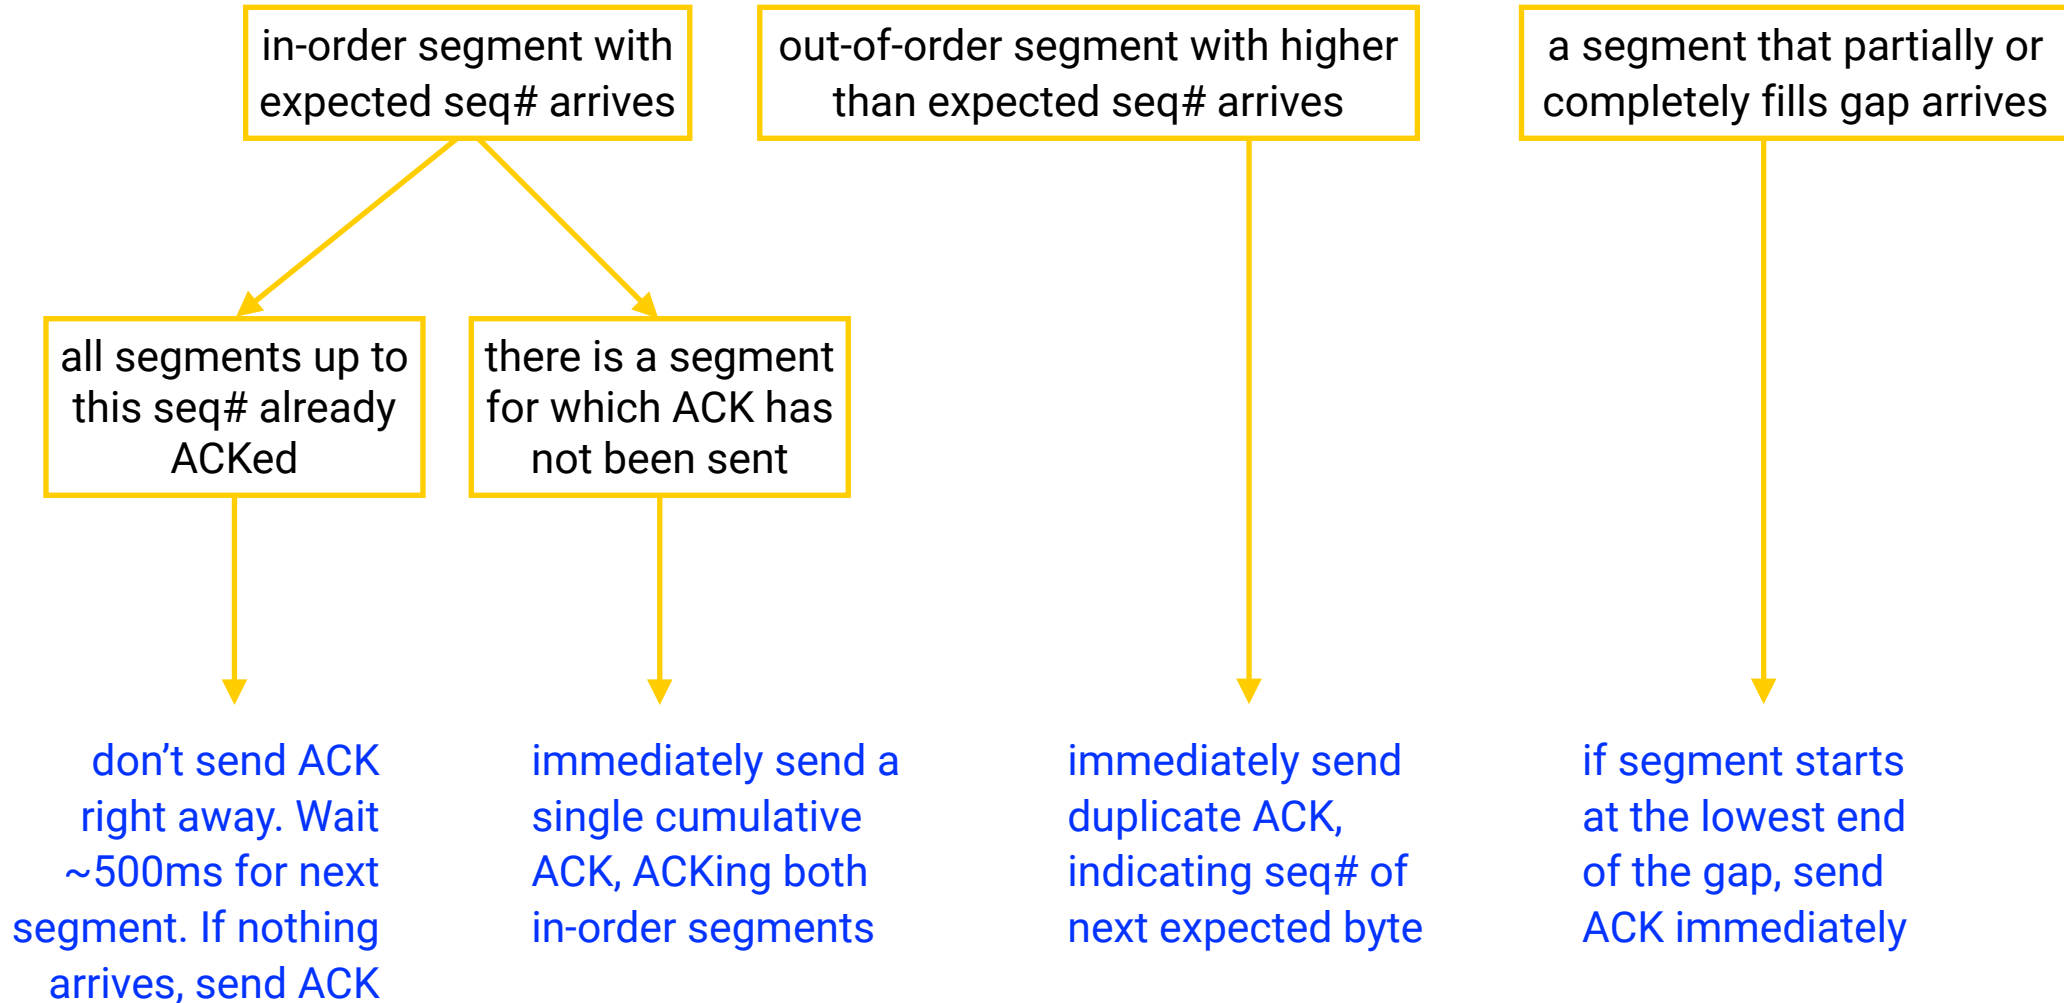


cumulative ACK covers
for earlier lost ACK



premature timeout

ACK Generation by TCP Receiver (RFC 5681)

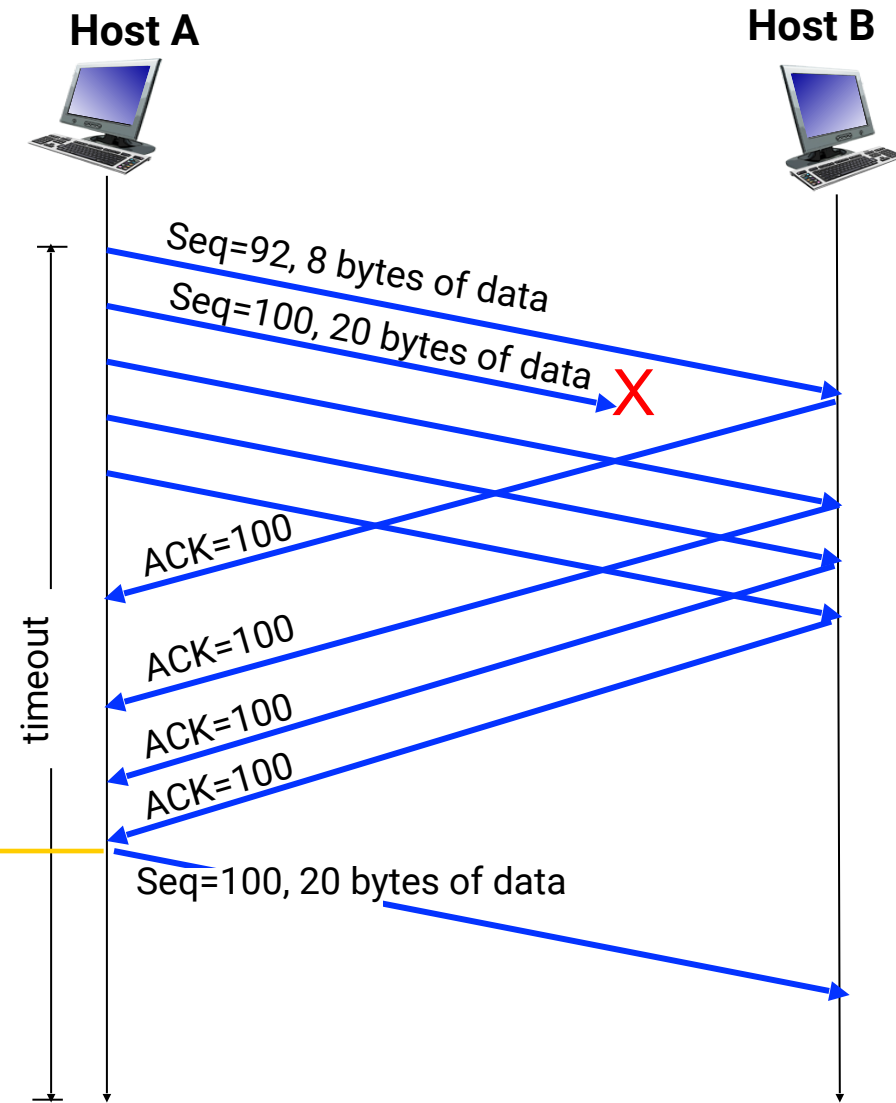


TCP fast retransmit

if sender receives 3 ACKs for same data ("triple duplicate ACKs"), it is likely that unacknowledged segment is lost, so don't wait for timeout, instead resend that segment now



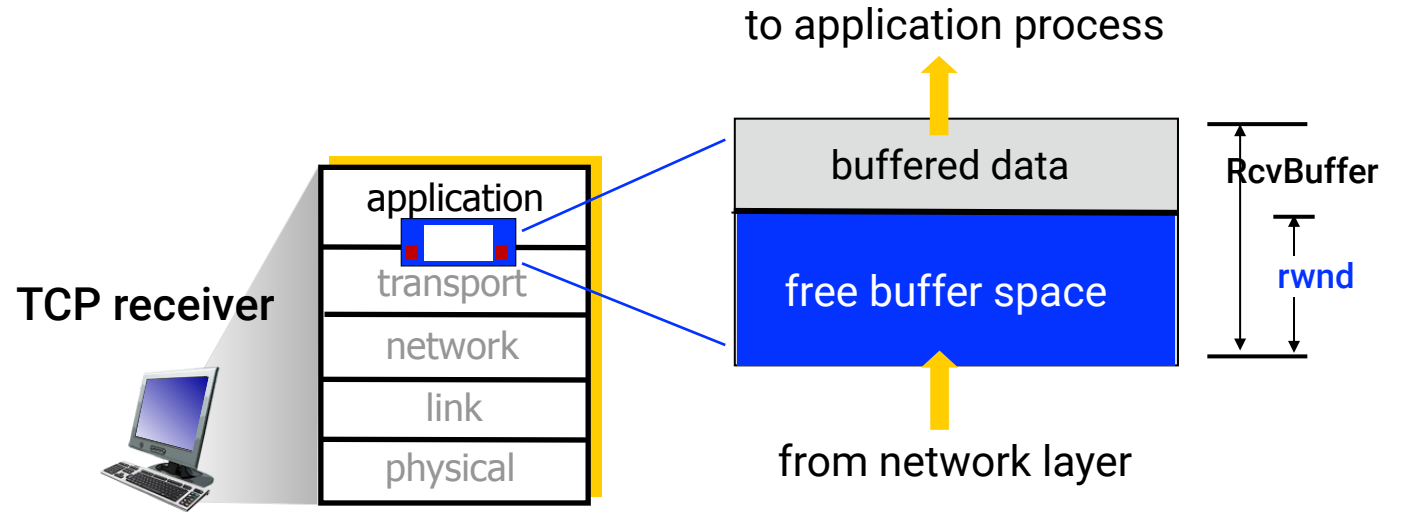
Receipt of triple duplicate ACKs indicates 3 segments received after a missing segment, so lost segment is likely. Retransmit!



TCP Flow Control

Key idea

let the receiver control the sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



- TCP receiver advertises its free buffer space in **rwnd** field in TCP header
- rwnd is typically set to 4kB, while its full range is 0 to 64kB (16-bit field)
- managed internally by the TCP/IP stack, and could be modified via socket options()
- sender limits amount of unacknowledged, in-flight data to receiver's **rwnd**, thereby guaranteeing that receiver won't experience buffer overflow

TCP Congestion Control


Key idea: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event

Additive Increase

increase sending rate by 1 maximum segment size every RTT until loss detected

Multiplicative Decrease

cut sending rate in half at each loss event (e.g., triple dup ACKs)




AIMD

- sawtooth behavior: probing for bandwidth
- a distributed, asynchronous algorithm
- shown to optimize network-wide flow rates

Classical TCP Implementation

sender sequence number space



send cwnd bytes, wait RTT for ACKS,
then send more bytes

TCP sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

cwnd is dynamically adjusted in response
to observed network congestion events

$$\text{TCP rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$


Two Phases: Slow Start and Congestion Avoidance

Slow Start: when a connection begins, increase sending rate exponentially until the first loss event

- start with `cwnd` = 1 MSS
- double `cwnd` every RTT i.e., increment `cwnd` for every ACK received

Congestion Avoidance: switch from exponential increase to linear increase when the connection hits first timeout

- set `ss-threshold` = `cwnd`/2
- switch to additive increase anytime `cwnd` reaches this level in the future



Spot Quiz (ICON)