

IOWA

CS5630

---

# Foundation (3): Compute in the Cloud

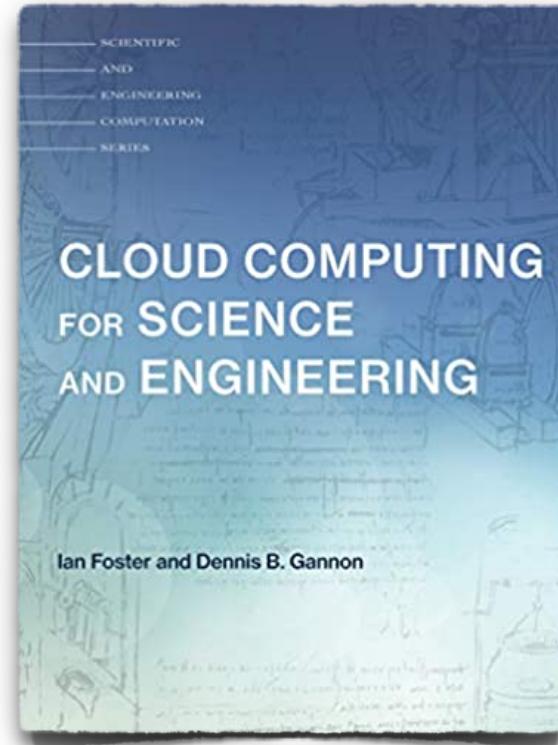
**Prof. Supreeth Shastri**  
*Computer Science*  
*The University of Iowa*

# Lecture goals

---

*Technical introduction to computing  
resources on the cloud*

- **Infrastructure as a Service (IaaS)**
- **Platform as a Service (PaaS)**
- **Software as a Service (SaaS)**
- **Other emerging models**

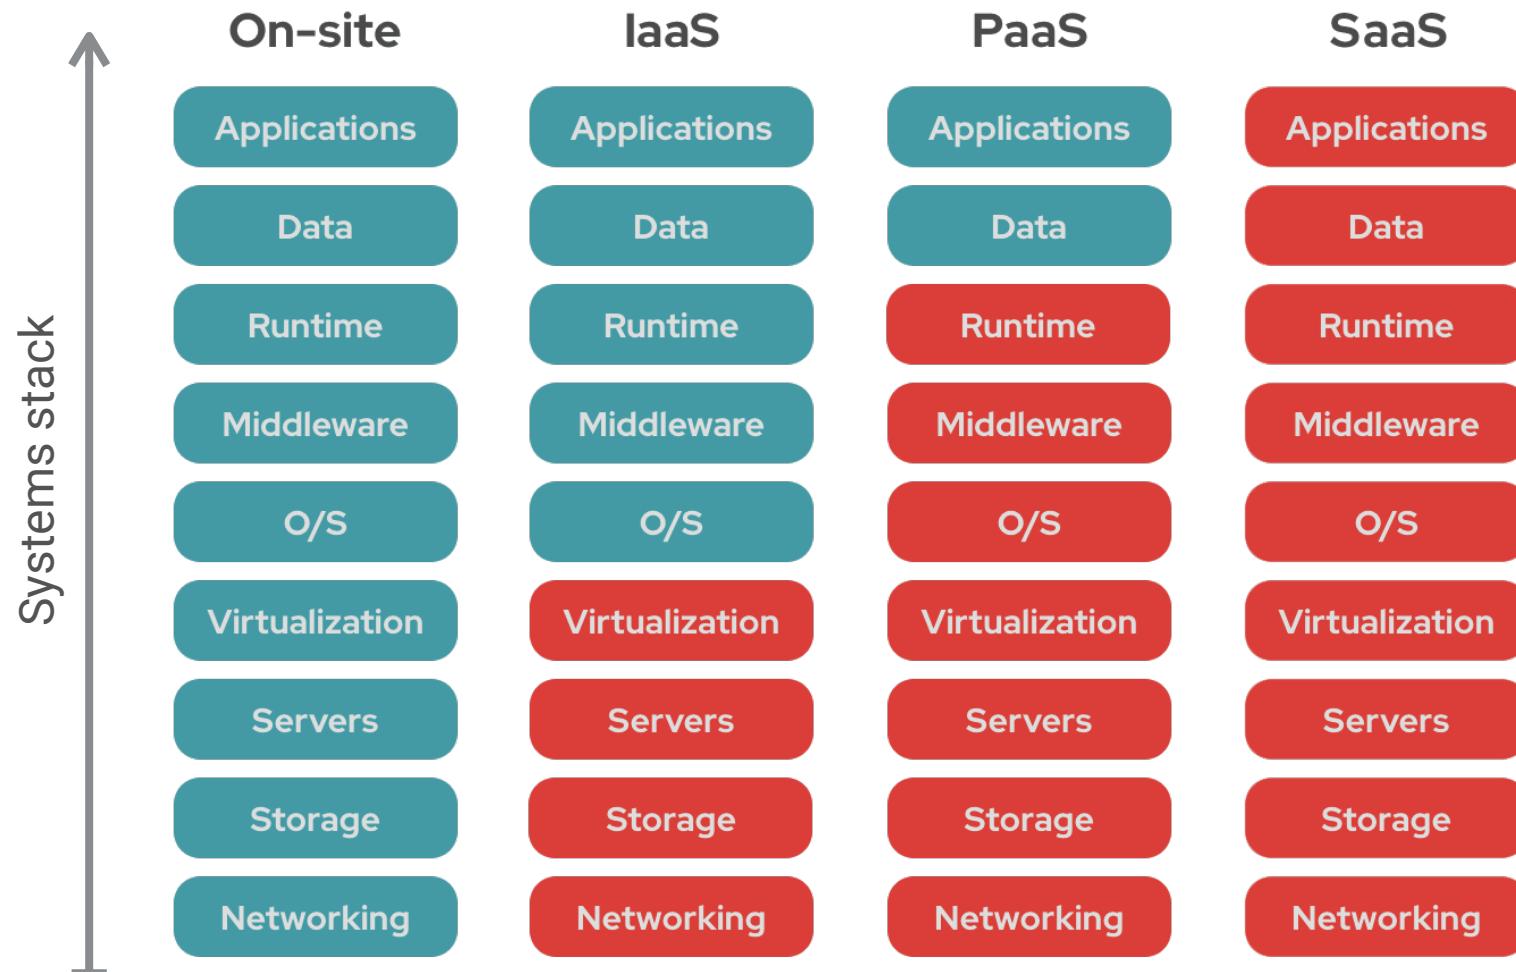


Chapters 4, 5



Google Cloud

# Visualizing compute in the cloud



■ You manage

■ Service provider manages

Image courtesy: Redhat

# IaaS Offerings are Redefining the Compute Server

---



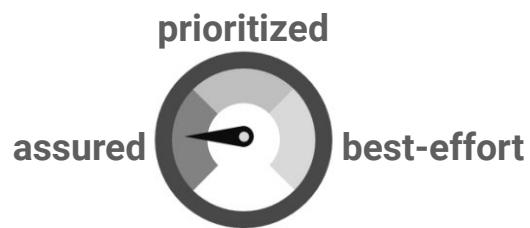
The Cloud Era

## Static characteristics

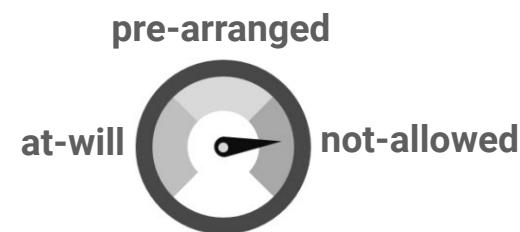
*Price, performance, and failure characteristics that do not change in an application's lifetime*

“A rose by any other name *may not* smell as sweet”

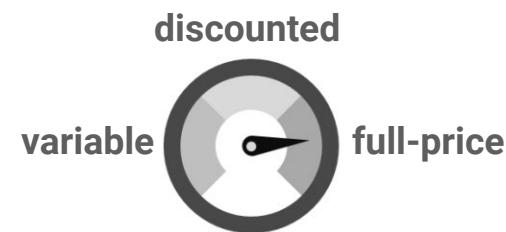
*All service-level characteristics are controlled in fine granularity*



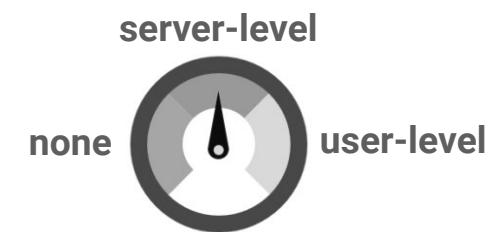
## Availability



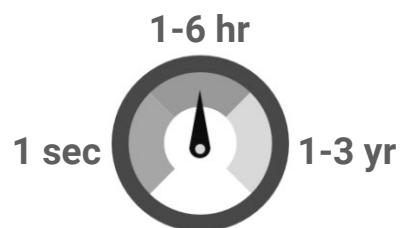
## Revocability



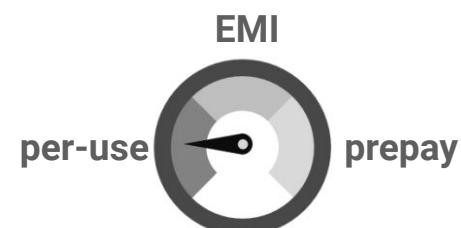
## Pricing Model



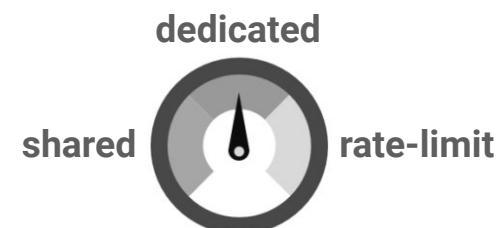
## Isolation



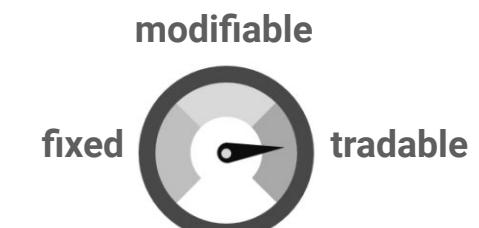
## Time Commitment



## Billing Frequency



## Performance Guarantees



## Terms of Contract



## Availability

## Revocability

## Pricing Model

## Isolation

**Spot Server + 15 other IaaS offerings on EC2**



## Time Commitment

## Billing Frequency

## Performance Guarantees

## Terms of Contract

# EC2 Contract Diversity

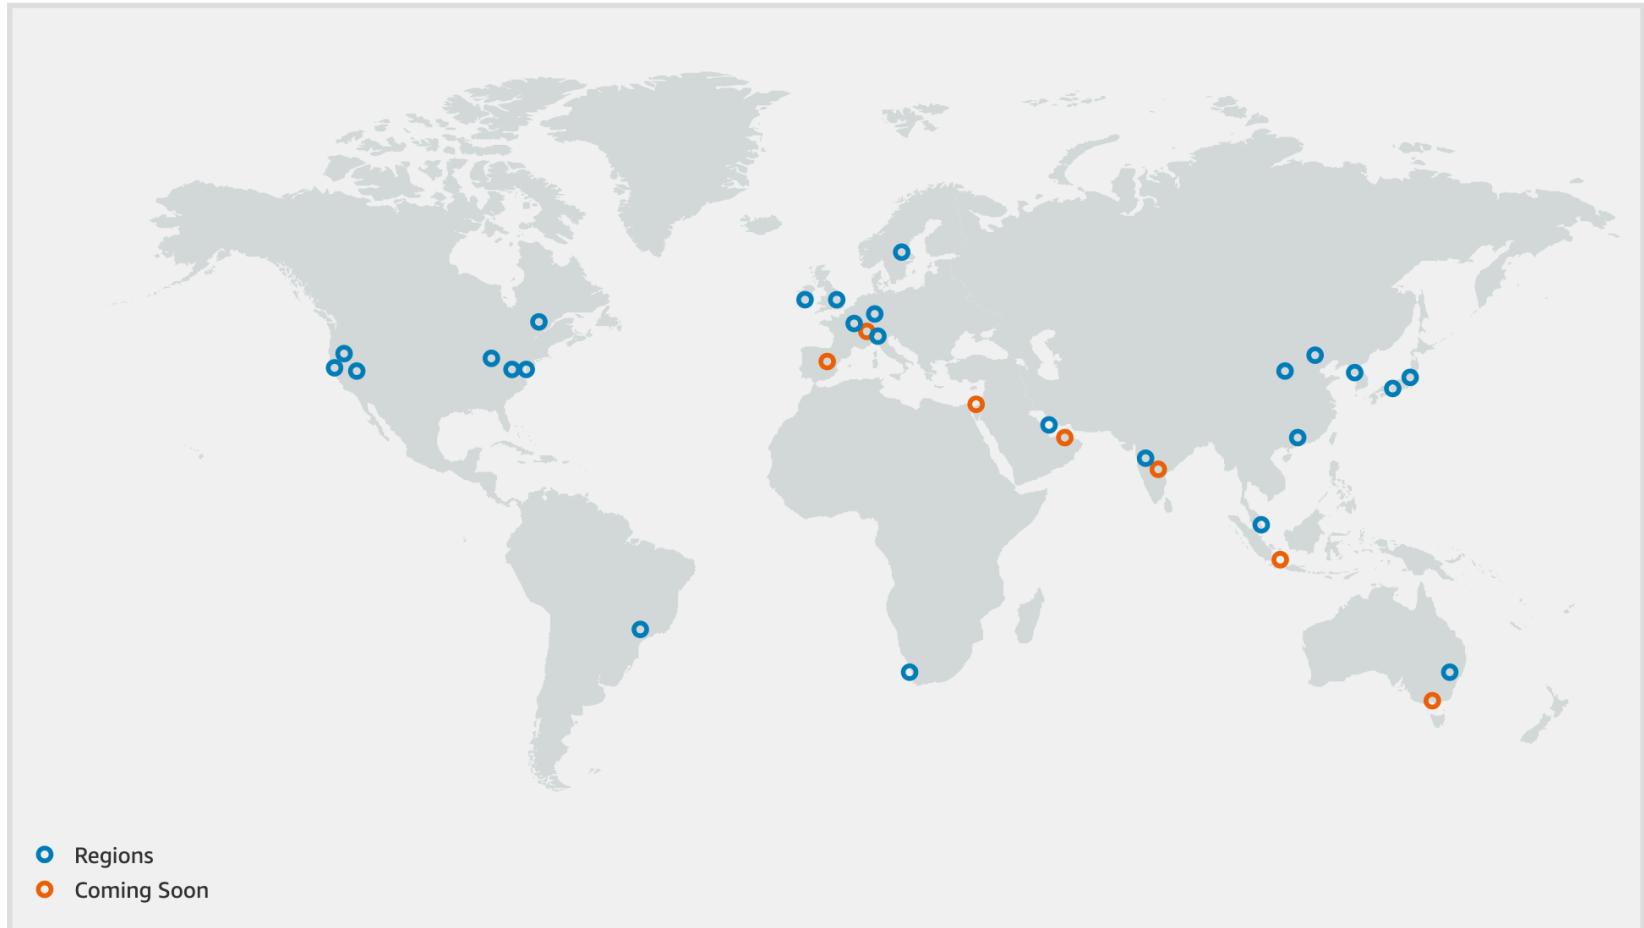
---

EC2 contract	Cost model	Performance	Availability	Revocation risk
On-demand	fixed cost per sec	low variability	high	no
Reserved	fixed cost per year	low variability	guaranteed	no
Spot	variable cost per sec	low variability	not guaranteed	yes (market driven)
Defined-duration	variable cost per hour	low variability	not guaranteed	yes (on expiry)
Burstable	fixed cost per sec	high variability	high	no
Dedicated	fixed cost per sec	no variability	high	no

# EC2 Location Diversity

---

- **25 geographic regions** (2x more than the next largest provider)
- **81 availability zones** (i.e., fully isolated datacenters)
- **230+ edge** locations/caches



# EC2 Hardware Spec

## Five categories

- ➡ General purpose
- ➡ Compute optimized
- ➡ Memory optimized
- ➡ Storage optimized
- ➡ Accelerated

**400**

instance types

	C6g	C6gn	C5	C5a	C5n	C4		
Model	vCPU	Memory (GiB)		Instance Storage (GiB)		Network Bandwidth (Gbps)***	EBS Bandwidth (Mbps)	
c5.large	2	4		EBS-Only		Up to 10	Up to 4,750	
c5.xlarge	4	8		EBS-Only		Up to 10	Up to 4,750	
c5.2xlarge	8	16		EBS-Only		Up to 10	Up to 4,750	
c5.4xlarge	16	32		EBS-Only		Up to 10	4,750	
c5.9xlarge	36	72		EBS-Only		10	9,500	
c5.12xlarge	48	96		EBS-Only		12	9,500	
c5.18xlarge	72	144		EBS-Only		25	19,000	
c5.24xlarge	96	192		EBS-Only		25	19,000	
c5.metal	96	192		EBS-Only		25	19,000	



*location / hardware*

**Contract diversity is omnipresent in the real world**

---

**Short-term  
rental**

---



**Assisted  
living**



**Long-term  
room share**



**Lease-to-own**

---



# Pricing the EC2 Servers

- Price is a function of hardware spec, location, and choice of contract
- To illustrate, let's consider two instance types in two regions under five different contracts
  - c5 series is compute-optimized and x1e series is memory optimized
  - us-east-1 is located in Virginia, whereas ap-northeast-3 is in Osaka

*prices decrease this way*

c5.24xlarge	us-east-1	ap-northeast-3
On-demand	4.080	5.136
Dedicated	3.366	4.237
Reserved (1y)	2.570	3.236
Defined-duration	2.244	2.825
Spot	1.555	1.596

*price diff. reduced*

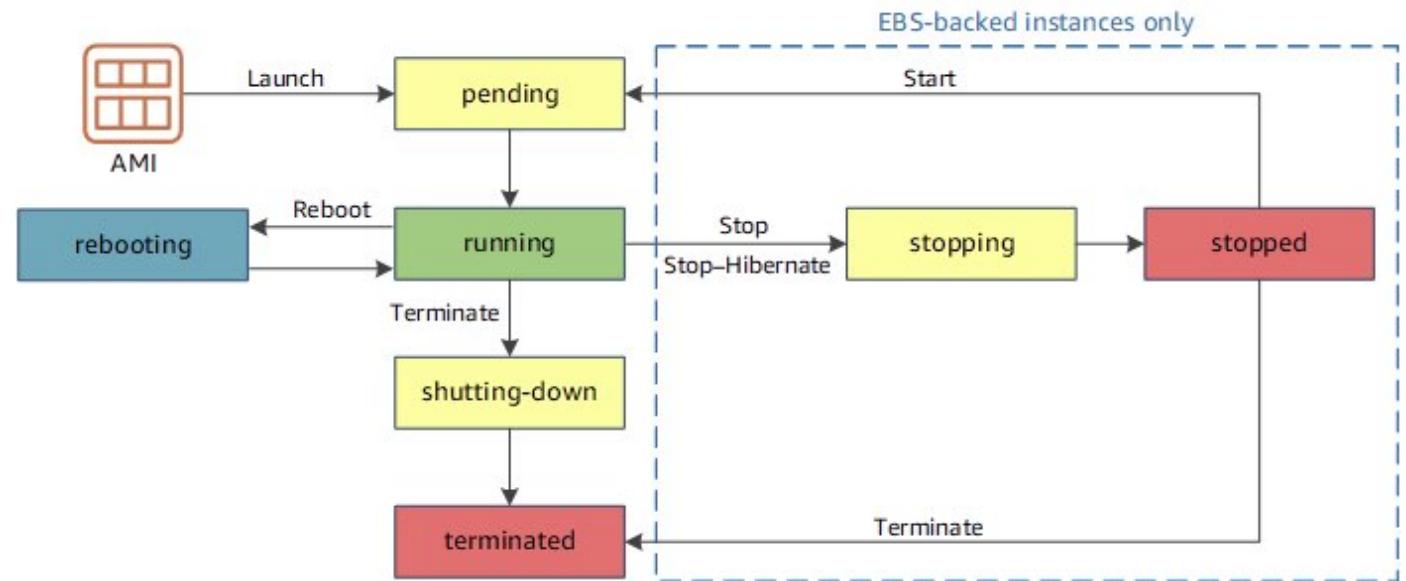
x1e.32xlarge	us-east-1	ap-northeast-3
On-demand	26.688	38.688
Dedicated	29.357	42.557
Reserved (1y)	16.446	23.839
Defined-duration	14.678	21.278
Spot	8.006	N/A

*unavailability*

*← price rise*

# **Working with VM on EC2**

# EC2 instance lifecycle



- ➡ This diagram should be familiar from your OS class (process lifecycle, VM lifecycle etc)
- ➡ All instances will be in one of the seven states; but billed only in the <running> state
- ➡ Hibernation (shown in blue box) is only possible when the root file system is backed by EBS
- ➡ Instance lifecycle can be managed via a web console, CLIs, or SDK-supported APIs
  - *For this exercise, we will use AWS CLIs*
  - *AWS CLI is an open-source tool that allows full management via command-line shell of Linux/MacOS/Windows*

## 1. Create a key pair

## 2. Set up a security group

## 3. Launch instance

## 4. Associate public IP

## 5. Connect to instance

## 6. Manage lifecycle

```
$ aws ec2 create-key-pair --key-name cs5630 --output text > cs5630.pem
```

-----BEGIN RSA PRIVATE KEY-----

```
EXAMPLEKEYCAQEAY7WZhaDsrA1W3mRlQtvhwyORRX8gnxgDafRt/gx42kWXsT4rXE/b5CpSgie/vBoU7jLxx92pNHoFnByP+Dc21eyyz6CvjTmWA0JwfWiW5/akH7i05dSrvC7dQkw2duV5QuUdE0QWZ/aNxMniGQE6XAgfwlnXBwrerrQo+ZWQeqiUwwMkuEbLeJFLhMCvYURpUMSC1oehm449ilx9X1FG50TCFeOzf18dqqCP6GzbPaIjiU19xx/azOR9V+tpUoZEL+wmXnzt3/nHPQ5xvD2OJH67km6SuPWoPzev/D8V+x4+bHthfsjR9Y7DvQFjfBVwHXigBdtZcU2/wei8D/HYwIDAQABaoIBAGZ1kaEvnrqu/uler7vgIn5m71n5Lkw4hJLAIW6tUT/fzvtcHK0SkbQCQXuriHmQ2MQyJX/0kn2NfjLV/ufGxbL1mb5qwMGUnEpJaZD6QSSs3kICLwWUYUiGfc0uiSbmJoap/GTLU0W5Mfcv36PaBUNy5p53V6G7hx2bahyWyJNfjLe4M86yd2YK3V2CmK+X/BOsShnJ36+hjrXPPWmV3N9zEmCdJja+K15DYmhm/tJWSD981oGk9TopEp7CkIfatEATyyZiVqoRq6k64iuM9JkA3OzdXzMqexXVJ1TLZVEH0E7bh1Y9d801ozRoQs/FiZNAX2iijCWyv0lpje73+kCgYEAm9mZtyhkHkFDpwrSM1APaL8oNAbbjwEy7Z5Mqfq1+lIp1YkriL0DbLx1vRAH+yHPRit2hHOjtUNZh4Axv+cpg09qbUI3+43eEy24B7G/Uh+GTfbjsXs0xQx/xp9otyVwc7hsQ5TA5PZb+mvkJ5OBEKzet9XcKwONBYELGhnEPe7cCgYEAm6Vgov6YHleHui9kHuwsayav0elc5zkxjF9nfHFJRRry21R1trw2Vdpn+9g481URrpzWVOEihvm+xTtmaZ1Sp//lkq75XDwnUWA8gkn603QE3fq2yN98BURsAKdJfJ5RL1HvGqvTe10HLYYXpJnEkHv+Unl2ajLivWUt5pbBrKbUCgYBjb0+OZk0sCcpZ29sbzjYjpIdxErySIyRX5gv2uNQwAjLdp9Pfn295yQ+BxMBXiIycWVQiwbh0Mo7yykABY7Ozd5wQewBQ4AdS1WSX4nGDtsiFxwi5sKuAAeOCbTosyls8w8fxoJ5Tz1sdoxNeGsArq6Wv/G16zQuAE9zK9vvwKBgF+09VI/1wJBirsDGz9whVWffPrTkJNvJzzYt69qezxlsjgFKshyWBhd4xHZtmCqpBP1AymEjr/T0lbxyARmXMnIOWIAAnNXMGB4KGSS11mzSVAoQ+fqr+cJ3d0dyP11jjjb0Ed/NY8frlNDxAVHE8BSkdsx2f6ELEyBKJSRr9snRAoGAMrTwYneXzvTskF/S5Fyu0iOegLdaNWUH38v/nDCgEpIXD5Hn3qAEcju1ijmbwlvtW+nY2jVhv7UGd8MjwUTNGItdb6nsYqM2asrnF3qSVrkAKKKYeGjkpUfvTrw0YFjXkfcrR/V+QFL5OndHAKJXjw7a4ejJLncTzmZSpYzwApc=-----END RSA PRIVATE KEY-----
```

1. Create a key pair
2. Set up a security group
3. Launch instance
4. Associate public IP
5. Connect to instance
6. Manage lifecycle

```
$ aws ec2 create-security-group --group-name mySG --description "My SG"
{
    "GroupId": "sg-903004f8"
}
```

1. Create a key pair

2. Set up a security group

3. Launch instance

4. Associate public IP

5. Connect to instance

6. Manage lifecycle

```
$ aws ec2 run-instances --image-id <ami-173d747e> --count 1 --instance-type t3.micro --key-name cs5630 --security-groups mySG
{
    "OwnerId": "123456789012",
    "ReservationId": "r-5875ca20",
    "Instances": [
        {
            "PublicDnsName": "ec2-192-0-2-1.us-west-2.compute.amazonaws.com",
            "Platform": "linux",
            "State": {
                "Code": 0,
                "Name": "pending"
            },
            "LaunchTime": "2021-09-05T02:42:39.000Z",
            "InstanceId": "i-5203422c",
            "ImageId": "ami-173d747e",
            "KeyName": "cs5630",
            "InstanceType": "t3.micro",
            "Placement": {
                "AvailabilityZone": "us-west-2b"
            },
            "BlockDeviceMappings": [
                {
                    "DeviceName": "/dev/sda1",
                    "Ebs": {
                        "Status": "attached",
                        "DeleteOnTermination": true,
                        "VolumeId": "vol-877166c8",
                    }
                }
            ],
            "RootDeviceName": "/dev/sda1",
            "RootDeviceType": "ebs",
        }
    ]
}
```

1. Create a key pair
2. Set up a security group
3. Launch instance
4. Associate public IP
5. Connect to instance
6. Manage lifecycle

```
$ aws ec2 associate-address --instance-id i-5203422c --public-ip 198.51.100.0  
  
$ aws ec2 authorize-security-group-ingress --group-name mySG --protocol tcp --  
port 22 --cidr 198.51.100.0
```

1. Create a key pair
2. Set up a security group
3. Launch instance
4. Associate public IP
5. Connect to instance
6. Manage lifecycle

```
$ ssh -i cs5630.pem ec2-user@198.51.100.0  
$ scp -i cs5630.pem <src-file> ec2-user@198.51.100.0:<dest-path>
```

1. Create a key pair
2. Set up a security group
3. Launch instance
4. Associate public IP
5. Connect to instance
6. Manage lifecycle

```
$ aws ec2 stop-instances --instance-ids i-5203422c
{
    "StoppingInstances": [
        {
            "InstanceId": "i-5203422c",
            "CurrentState": {
                "Code": 64,
                "Name": "stopping"
            },
            "PreviousState": {
                "Code": 16,
                "Name": "running"
            }
        }
    ]
}

$ aws ec2 start-instances --instance-ids i-5203422c
$ aws ec2 reboot-instances --instance-ids i-5203422c
$ aws ec2 terminate-instances --instance-ids i-5203422c
```

# **Spot Quiz (ICON)**